

mlflow

Platform for Machine Learning Lifecycle

Jules S. Damji
@2twitme



Data Science for Social Good Summit

by Iberian DSSGx organizations, chapters of the DSSG Foundation

ONLINE EVENT

6-7, 13-14, 20-21 OCTOBER 2020



GitHub and Slides URL

GitHub: <https://github.com/dmatrix/ds4g-workshop>

Slides: Under the directory *slides*

Agenda

- Concepts and Motivation for MLflow
- MLflow Components
- MLflow Tutorial
- Q & A

About Me



Jules S. Damji
Senior Developer Advocate @
Databricks

Joined Databricks in 2016
20+ Developer at software companies:
Sun, Netscape, VeriSign, @Home, Scalix,
LoudCloud/Opsware, Hortonworks etc



databricks

**Unified data analytics platform for
data science, data engineering, and business analytics to solve tough
data problems**

Global company with 5,000 customers and 450+ partners

Original creators of popular data and machine learning open source projects



Machine Learning Development is Complex

Traditional Software vs. Machine Learning

Traditional Software

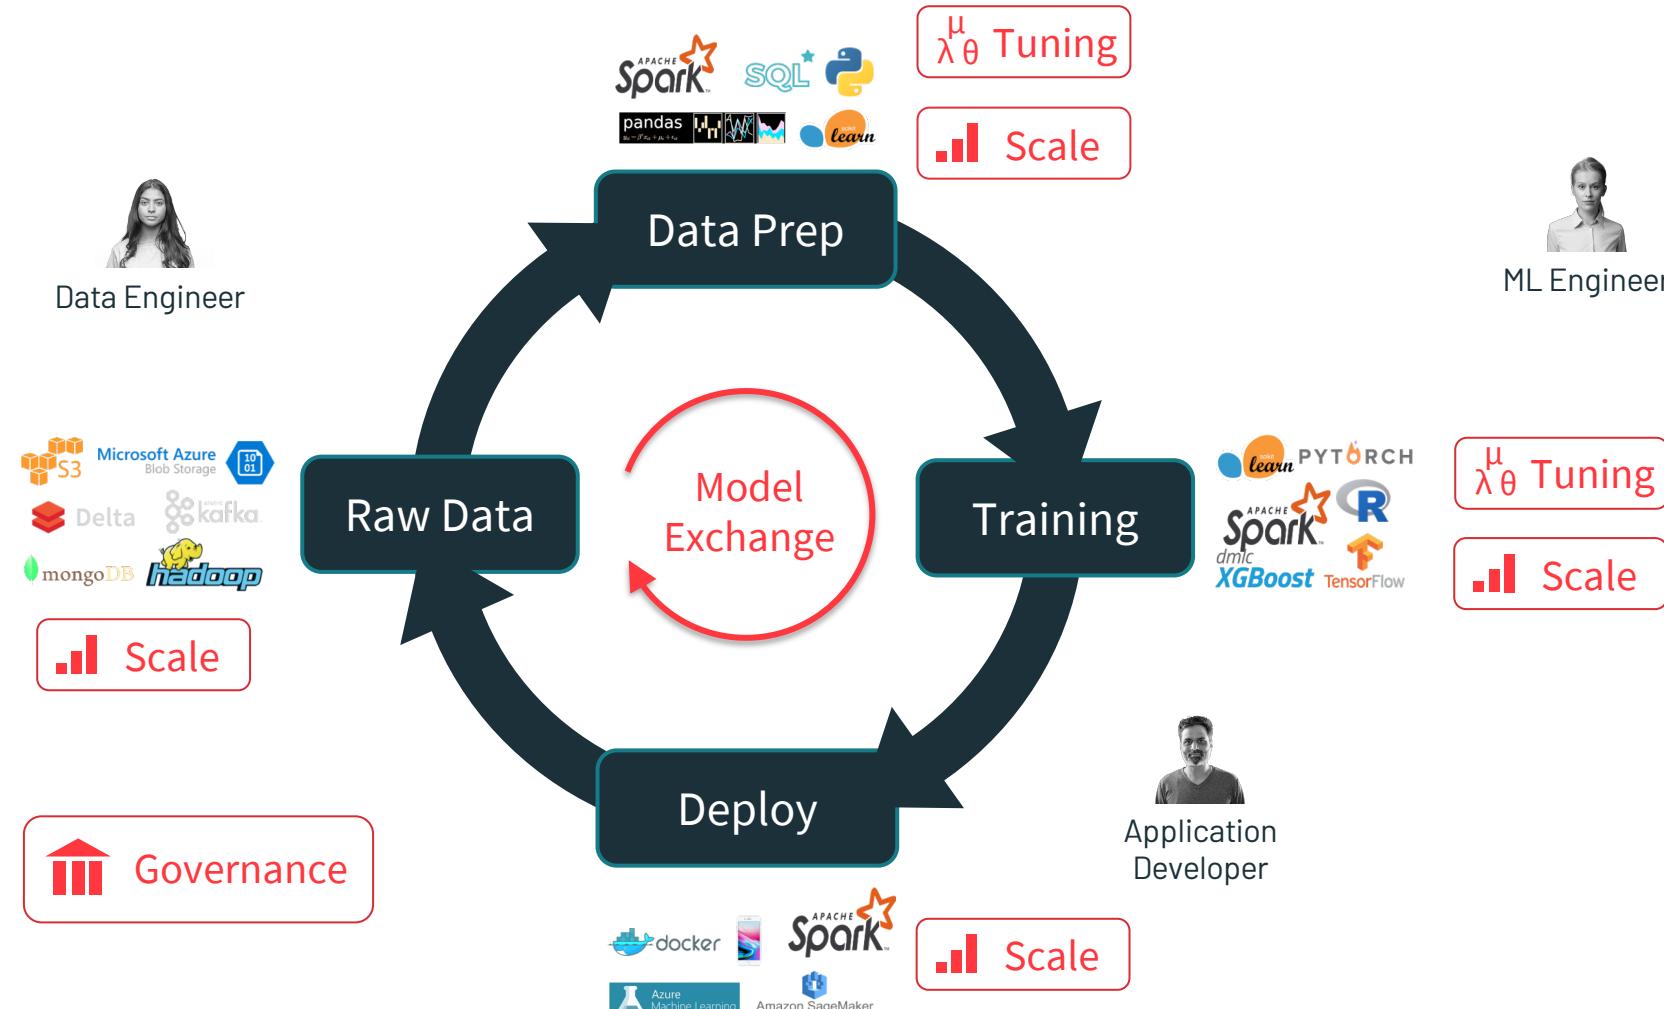
- **Goal:** Meet a functional specification
- Quality depends only on code
- Typically pick one software stack w/ fewer libraries and tools
- Limited deployment environments

Machine Learning

- **Goal:** Optimize metric (e.g., accuracy). Constantly experiment to improve it
- Quality depends on input data and tuning parameters
- Over time data changes; models drift...
- Compare + combine many libraries, model
- Diverse deployment environments



Machine Learning Lifecycle



Introducing **mlflow**

Open machine learning platform

Works with popular ML library & language

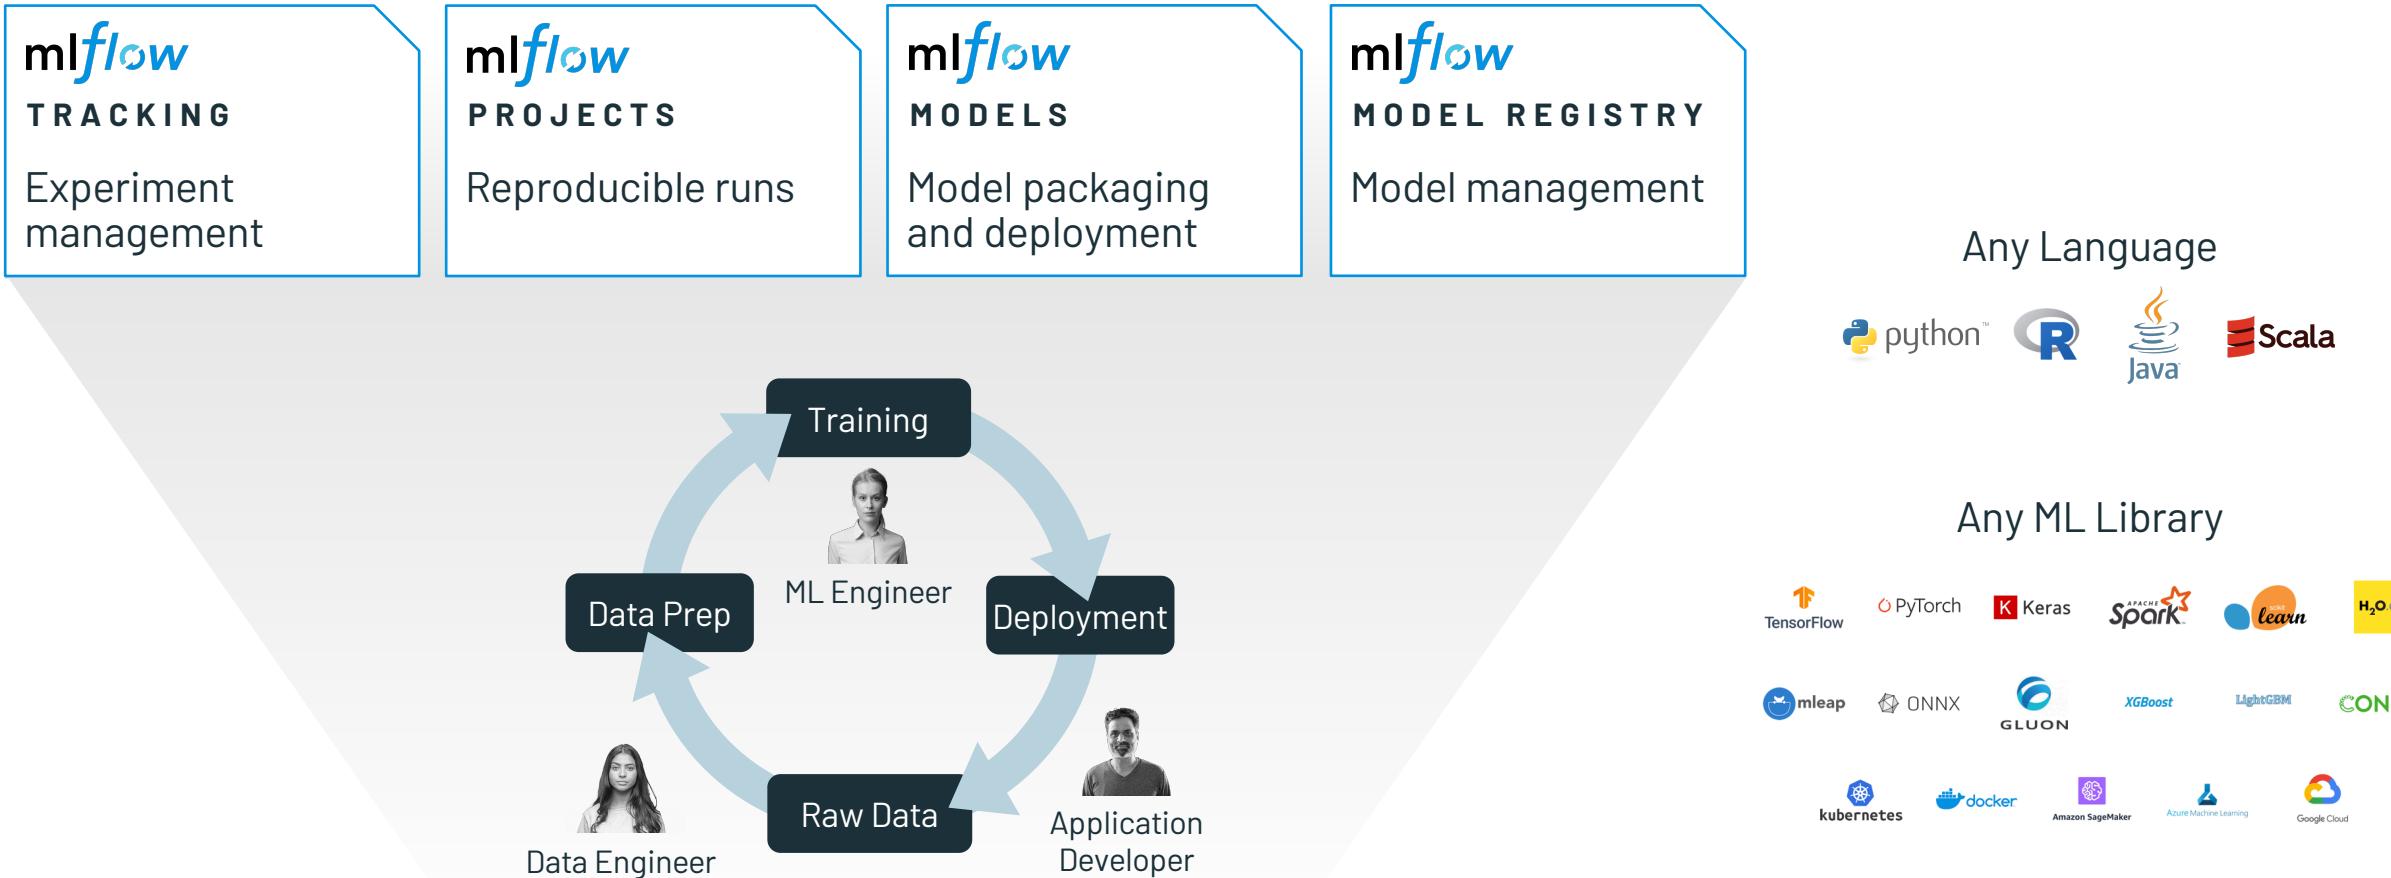
Runs the same way anywhere (e.g., any cloud or locally)

Designed to be useful for 1 or 1000+ person orgs

Simple. Modular. Easy-to-use.

Offers positive developer experience to get started!

mlflow: An Open Source ML Platform



MLflow Components

mlflow

Tracking

Record and query experiments: code, data, config, and results

mlflow

Projects

Package data science code in a format that enables reproducible runs on many platform

mlflow

Models

Deploy machine learning models in diverse serving environments

new

mlflow

Model Registry

Store, annotate and manage models in a central repository

[databricks.com
/mlflow](https://databricks.com/mlflow)



mlflow.org



github.com/mlflow



twitter.com/MLflow

Key Concepts in MLflow Tracking

Parameters: key-value inputs to your code

Metrics: numeric values (can update over time)

Tags and Notes: information about a run

Artifacts: files, data, and models

Source: what code ran?

Version: what of the code?

Run: an instance of code that runs by MLflow

Experiment: {Run, ... Run}

Model Development without MLflow Tracking

```
data    = load_text(file_name=file)
ngrams = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)

print("For n=%d, lr=%f: accuracy=%f"
      % (n, lr, score))

pickle.dump(model, open("model.pkl"))
```

```
For n=2, lr=0.1: accuracy=0.71
For n=2, lr=0.2: accuracy=0.79
For n=2, lr=0.5: accuracy=0.83
For n=2, lr=0.9: accuracy=0.79
For n=3, lr=0.1: accuracy=0.83
For n=3, lr=0.2: accuracy=0.82
For n=4, lr=0.5: accuracy=0.75
...
```

What version of
my code was this
result from?

Model Development with MLflow is Simple!

```
import mlflow
data    = load_text(file_name=file)
ngrams = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)
with mlflow.start_run():
    mlflow.log_param("data_file", file)
    mlflow.log_param("n", n)
    mlflow.log_param("learn_rate", lr)
    mlflow.log_metric("score", score)
    mlflow.sklearn.log_model(model)
```

\$ mlflow ui

The screenshot shows the MLflow UI interface. At the top, there's a search bar with the query "metrics.rmse < 1 and params.model = 'tree'". Below the search bar, there are sections for "Search Params" and "Filter Metrics". A table below lists 36 matching runs, each with a checkbox, Date, User, Source, Version, Parameters (alpha, l1_ratio), and Metrics (rmse, r2). The table has columns for Date, User, Source, Version, Parameters, and Metrics.

Date	User	Source	Version	Parameters	Metrics
2018-07-19 03:26:53	root	azure-demo1	0.01	0.55 0.596 0.25	0.596 0.25 0.762
2018-07-19 03:26:39	root	azure-demo	0.01	0.55	0.596 0.25
2018-07-19 03:26:14	root	azure-demo	0.01	0.55	0.596 0.25
2018-07-19 03:25:51	root	azure-demo	0.01	0.75	0.597 0.25
2018-07-19 03:25:42	root	azure-demo	0.01	0.04	0.591 0.256
2018-07-18 02:09:54	root	azure-demo	0.01	1.0	0.597 0.249
2018-07-18 02:09:29	root	azure-demo	0.01	0.75	0.597 0.25
2018-07-18 02:08:52	root	azure-demo	0.01	0.01	0.591 0.257
2018-07-17 08:13:37	root	azure-demo	0.01	0.01	0.591 0.257
2018-07-17 08:13:34	root	azure-demo	0.01	1.0	0.597 0.249
2018-07-17 08:13:30	root	azure-demo	0.01	0.75	0.597 0.25
2018-07-17 08:13:27	root	azure-demo	0.01	0.01	0.591 0.257
2018-07-17 08:08:05	root	azure-demo	0.01	0.01	0.591 0.257

Track parameters, metrics, artifacts, output files & code version

Model Development with MLflow is Simple!

mlflow.<flavor>.autolog() APIs

```
import mlflow
import mlflow.keras

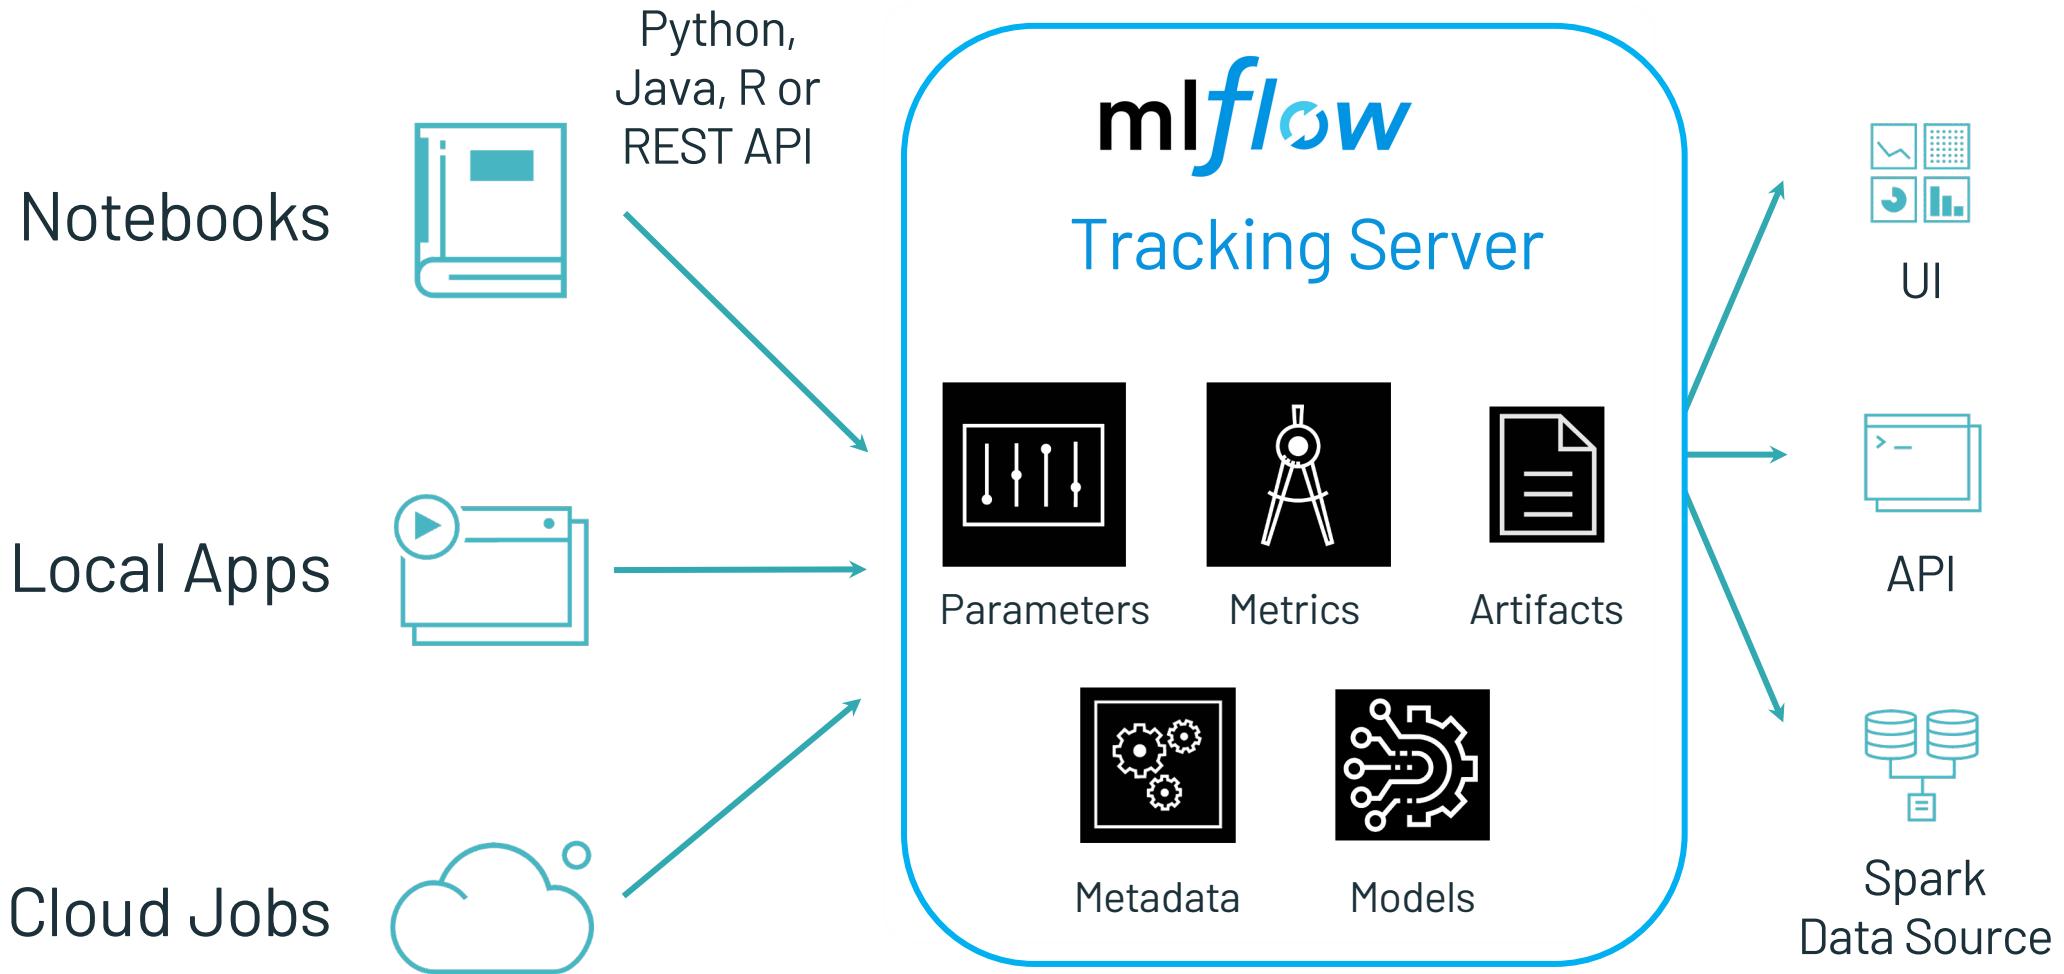
# Build, compile, enable autologging, and
# train your model
keras_model = ...
keras_model.compile(optimizer="rmsprop",
                     loss="mse",
                     metrics=["accuracy"])
# autolog your metrics, parameters, and
# model
mlflow.keras.autolog()
results = keras_model.fit( x_train, y_train,
                           epochs=20, batch_size=128,
                           validation_data=(x_val, y_val))
```

\$ mlflow ui

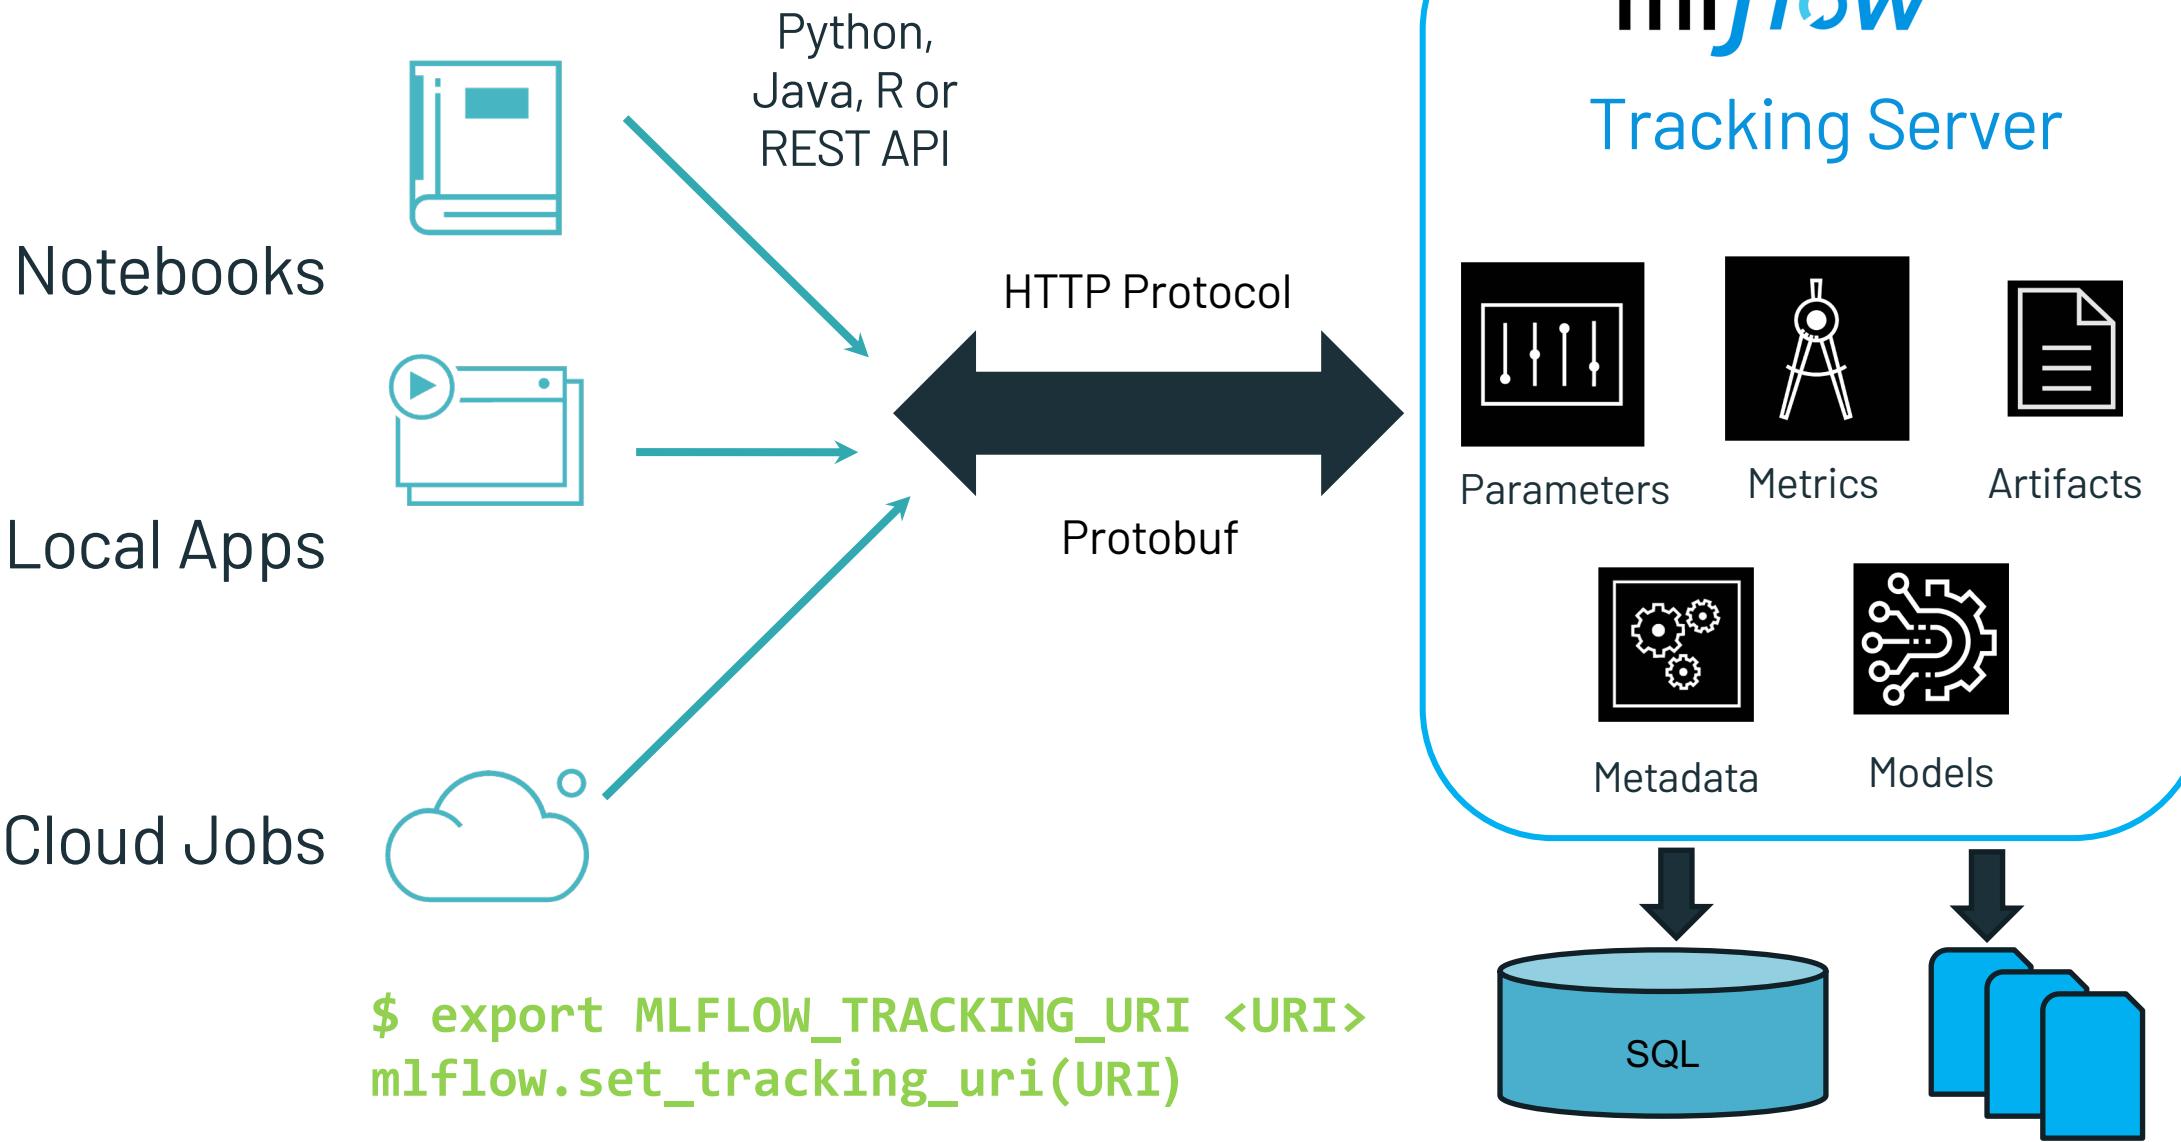
Date	User	Source	Version	Parameters	Metrics
2018-07-19 03:26:53	root	azure-demo1	0.01	0.55	0.596 0.25 0.762
2018-07-19 03:26:39	root	azure-demo	0.01	0.55	0.596 0.25 0.762
2018-07-19 03:26:14	root	azure-demo	0.01	0.55	0.596 0.25 0.762
2018-07-19 03:25:51	root	azure-demo	0.01	0.75	0.597 0.25 0.762
2018-07-19 03:25:42	root	azure-demo	0.01	0.04	0.591 0.256 0.759
2018-07-18 02:09:54	root	azure-demo	0.01	1.0	0.597 0.249 0.762
2018-07-18 02:09:29	root	azure-demo	0.01	0.75	0.597 0.25 0.762
2018-07-18 02:08:52	root	azure-demo	0.01	0.01	0.591 0.257 0.759
2018-07-17 08:13:37	root	azure-demo	0.01	0.01	0.591 0.257 0.759
2018-07-17 08:13:34	root	azure-demo	0.01	1.0	0.597 0.249 0.762
2018-07-17 08:13:30	root	azure-demo	0.01	0.75	0.597 0.25 0.762
2018-07-17 08:13:27	root	azure-demo	0.01	0.01	0.591 0.257 0.759
2018-07-17 08:08:05	root	azure-demo	0.01	0.01	0.591 0.257 0.759

Track parameters, metrics, artifacts, output files & code version

MLflow Tracking Server

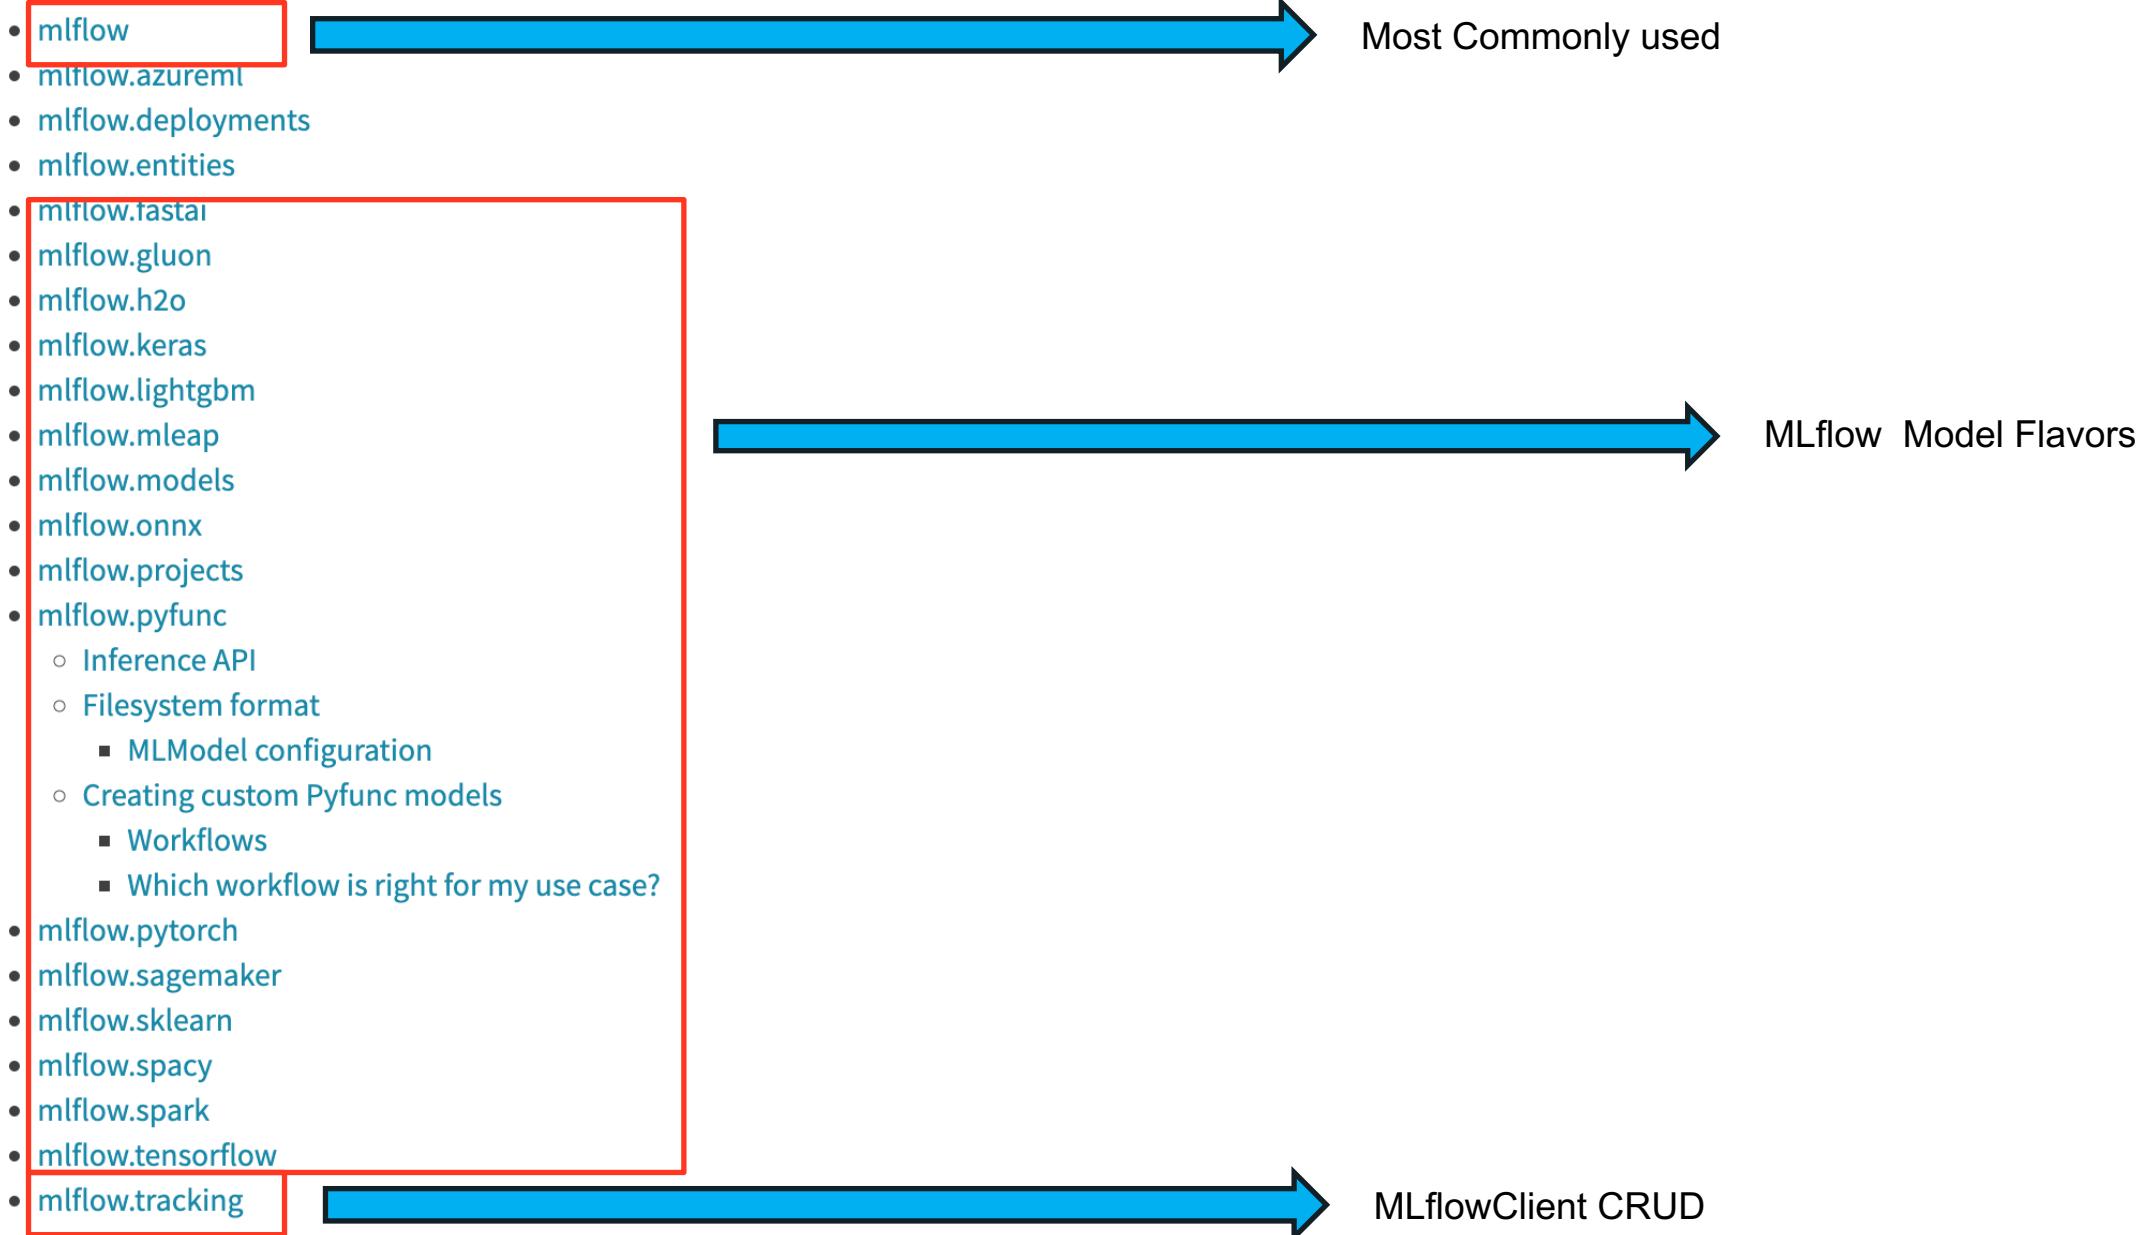


MLflow Tracking Server



Python API

The MLflow Python API is organized into the following modules. The most common functions are exposed in the `mlflow` module, so we recommend starting there.



Built-In Model Flavors

MLflow provides several standard flavors that might be useful in your applications. Specifically, many of its deployment tools support these flavors, so you can export your own model in one of these flavors to benefit from all these tools:

- [Python Function \(python_function\)](#)
- [R Function \(crate\)](#)
- [H2O \(h2o\)](#)
- [Keras \(keras\)](#)
- [MLeap \(mleap\)](#)
- [PyTorch \(pytorch\)](#)
- [Scikit-learn \(sklearn\)](#)
- [Spark MLlib \(spark\)](#)
- [TensorFlow \(tensorflow\)](#)
- [ONNX \(onnx\)](#)
- [MXNet Gluon \(gluon\)](#)
- [XGBoost \(xgboost\)](#)
- [LightGBM \(lightgbm\)](#)



mlflow.sklearn

The `mlflow.sklearn` module provides an API for logging and loading scikit-learn models. This module exports scikit-learn models with the following flavors:

Python (native) pickle format

This is the main flavor that can be loaded back into scikit-learn.

mlflow.pyfunc

Produced for use by generic pyfunc-based deployment tools and batch inference.

mlflow.sklearn.get_default_conda_env(include_cloudpickle=False) [source]

Returns

The default Conda environment for MLflow Models produced by calls to `save_model()` and `log_model()`.

mlflow.sklearn.load_model(model_uri) [source]

Load a scikit-learn model from a local file or a run.

Parameters

model_uri -

The location, in URI format, of the MLflow model, for example:

- `/Users/me/path/to/local/model`
- `relative/path/to/local/model`

MLFlow set of APIs

Fluent MLflow APIs

- Python
 - High-level operations for runs and experiments
 - Model Flavor APIs
- Java
 - MLflowContext
 - Experiments, runs, search, etc
- R
 - Experiments, runs, search etc

MLflowClient

- Low Level CRUD interface to experiments, runs, Entities, Model Registry, etc
- `import mlflow.tracking`
- `client = MLflowClient(**kwargs)`
- Metric, Param, Search etc

MLflow REST API

- Make REST calls to Tracking server with endpoints
- `https://<tracking_server>/api/..`
- `/2.0/mlflow/2.0/runs/log-metric`
- `/2.0/mlflow/runs/log-parameter`
- `/2.0/mlflow/runs/search`

MLflow Components

mlflow

Tracking

Record and query experiments: code, data, config, and results

mlflow

Projects

Package data science code in a format that enables reproducible runs on many platform

mlflow

Models

Deploy machine learning models in diverse serving environments

new

mlflow

Model Registry

Store, annotate and manage models in a central repository

[databricks.com
/mlflow](http://databricks.com/mlflow)



mlflow.org



github.com/mlflow



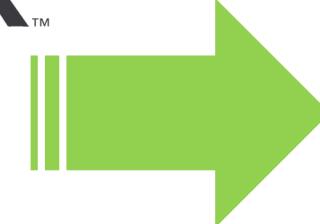
twitter.com/MLflow

MLflow Projects Motivation

Diverse set of tools



TensorFlow



Diverse set of environments

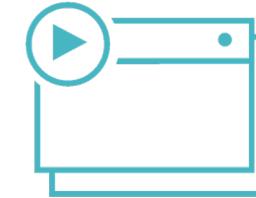
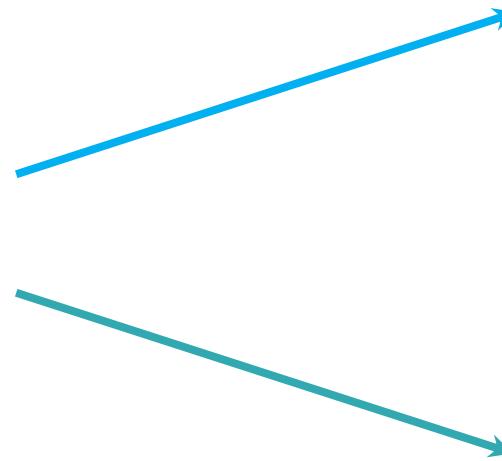
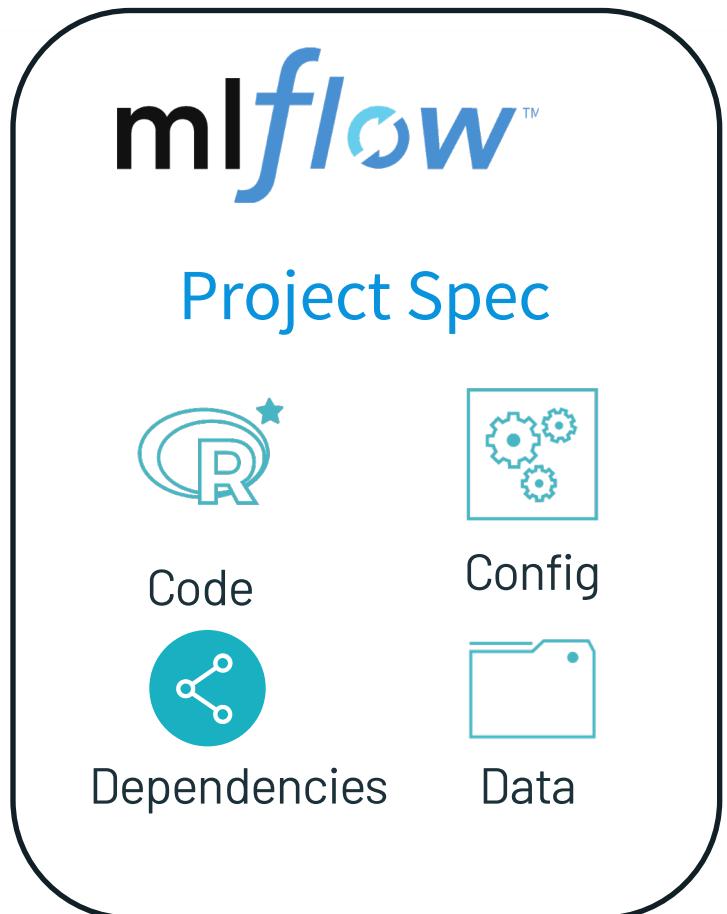


mlflow™
Projects

Package data science
code in a format that
enables reproducible runs
on any platform

Challenge: ML results difficult to reproduce

MLflow Projects



Local Execution



Remote Execution



1. Example MLflow Project File

```
my_project/
├── MLproject
|
└── main.py
    ├── conda.yaml
    └── model.py
...
...
```

```
name: tutorial

conda_env: conda.yaml

entry_points:
  main:
    parameters:
      batch_size: {type: int, default: 10}
      epochs: {type: int, default: 100}
    command: "python train_keras.py {batch_size} {epochs}
```

```
$ mlflow run git://<my_project>.git -P epochs=200
mlflow.run("git://<my_project>", parameters={..})

mlflow run . -e main -P epochs=200
```

2. Example Conda.yaml

```
my_project/
└── MLproject
    ├── conda.yaml
    ├── main.py
    └── model.py
    ...

```

```
channels:
- defaults
- conda-forge
dependencies:
- python=3.7.5
- pip
- pip:
  - mlflow
  - keras==2.3.1
  - tensorflow==2.0.0
name: mlflow-env
```

MLflow Projects

Packaging format for reproducible ML runs

- Any code folder or GitHub repository
- MLproject file with project configuration

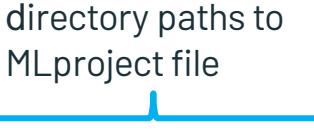
Defines dependencies for reproducibility

- Conda (+ R, Docker, ...) dependencies can be specified in MLproject
- Reproducible in (almost) any environment

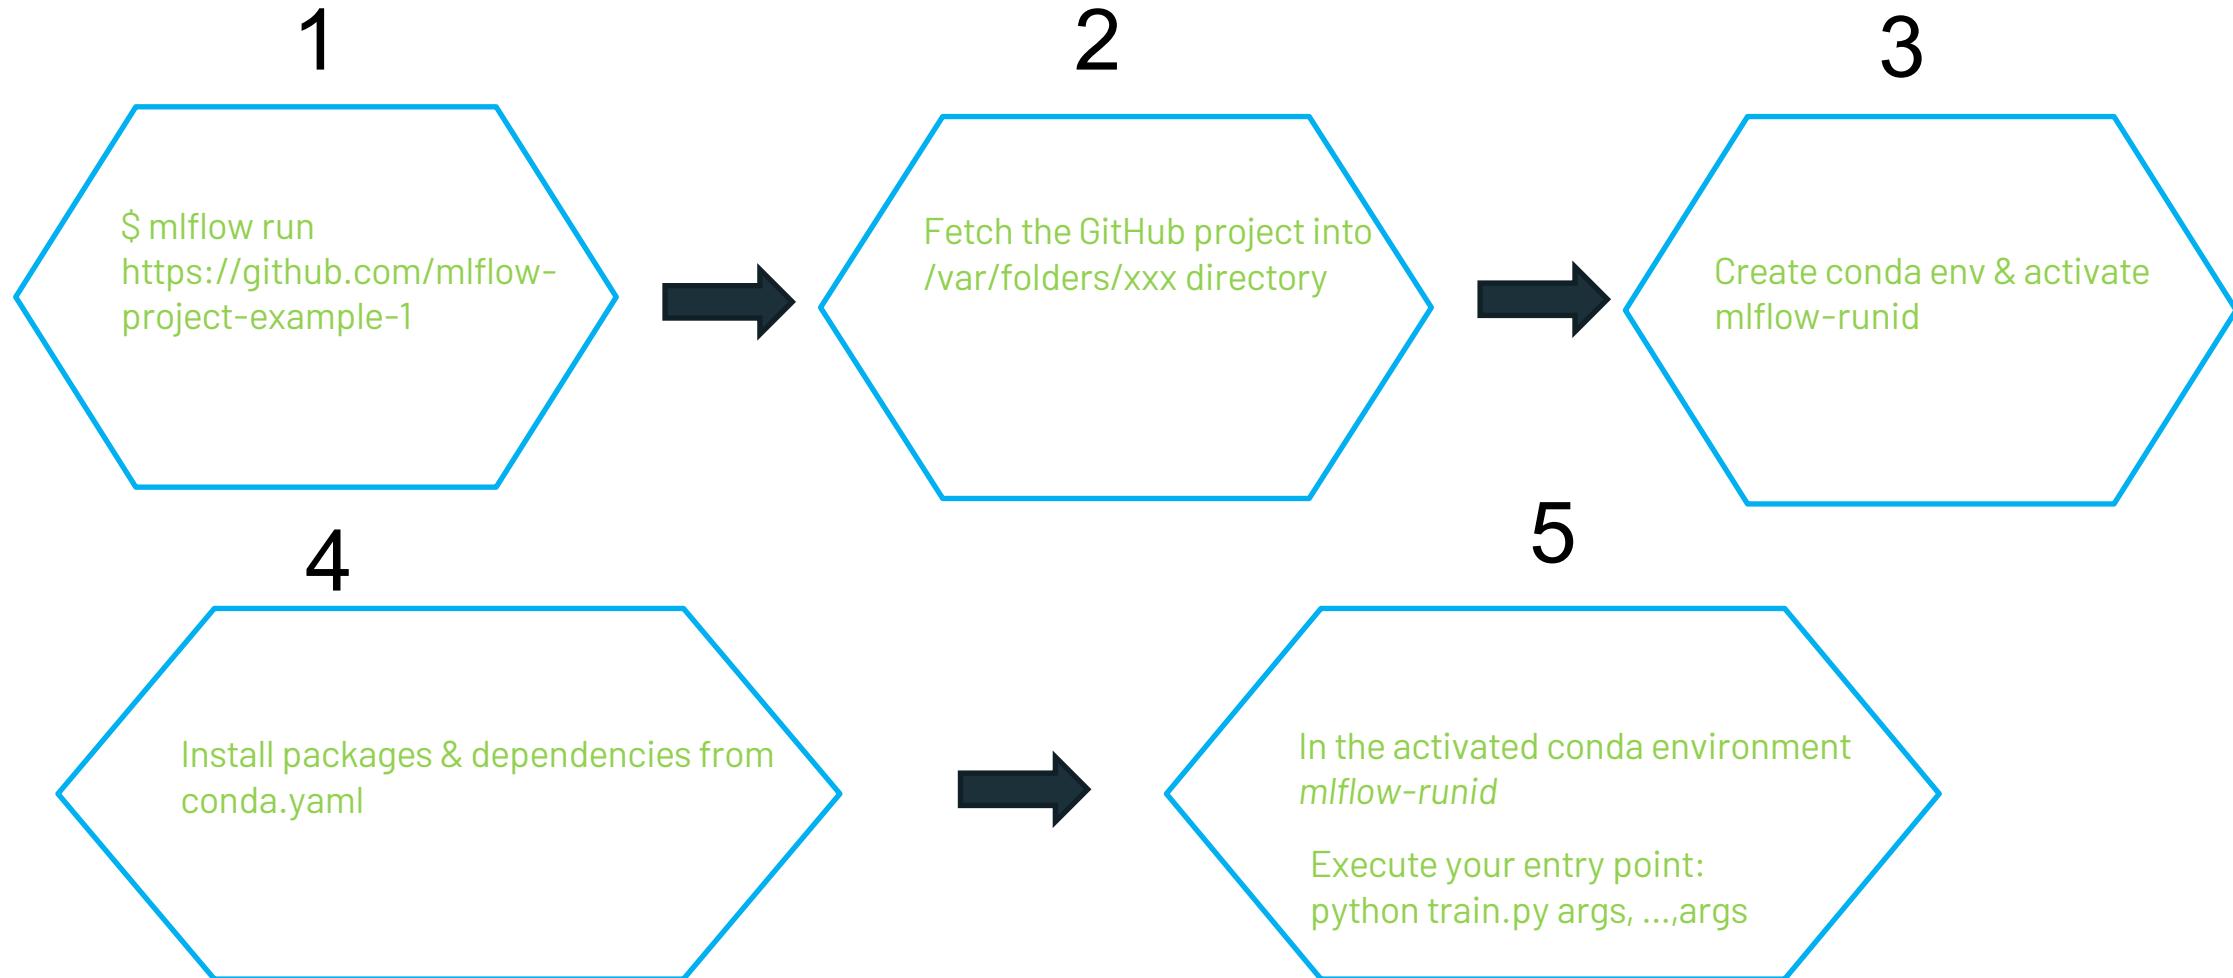
Execution API for running projects

- CLI / Python / R / Java
- Supports local and remote execution
 - mlflow run –help (CLI)
 - mlflow run <https://github.com/dmatrix/jsd-mlflow-examples.git#keras/imdbclassifier> (CLI)
 - mlflow.run (<project_uri>, parameters={}) or mlflow.projects.run(<project_uri>, parameters={}) (API)

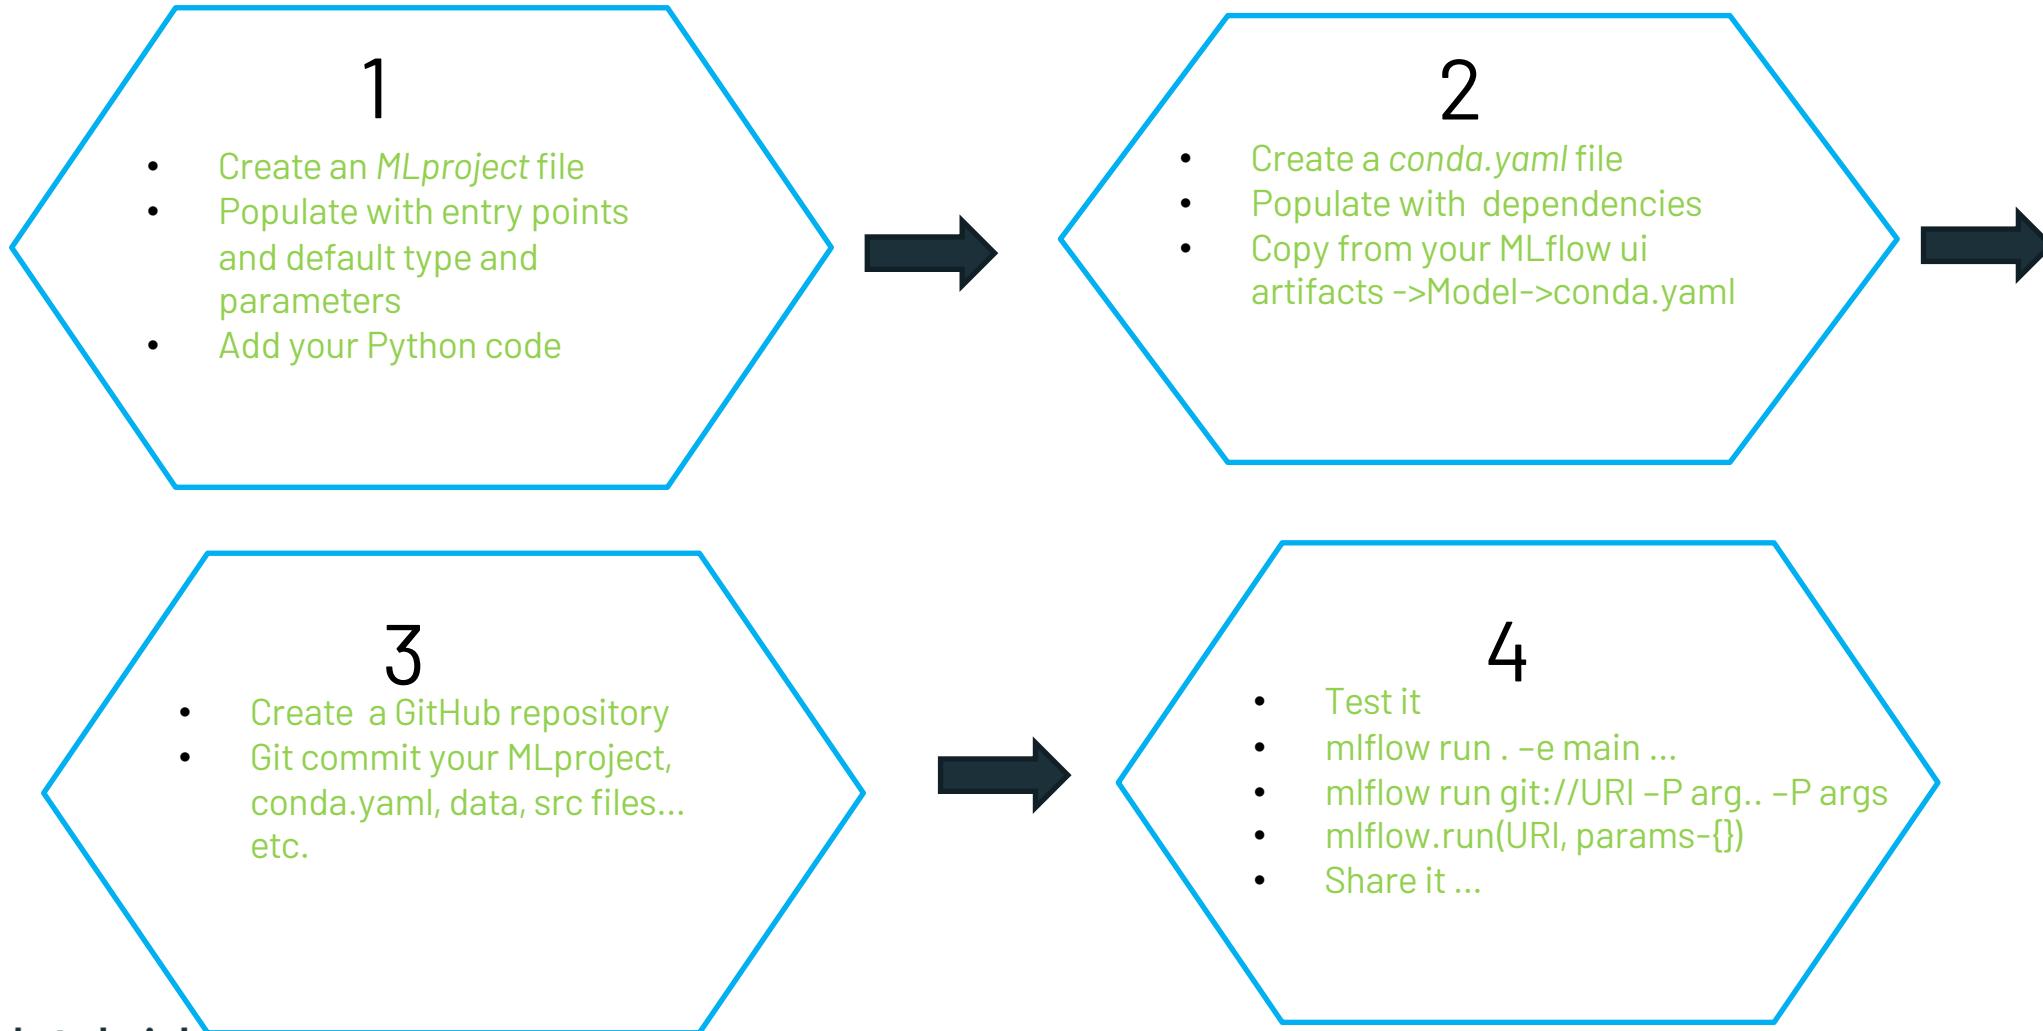
directory paths to
MLproject file



Anatomy of MLflow Project Execution



How to build an MLflow Project



MLflow Projects Usage

- Reproducibility and Experimentation ✓

- Reproduce experiments in a target environment
 - Experiment with different hyperparameters
 - Use of CLI and Programmatic interface

- Shareability ✓

- Use other people's models and experiment via a MLFlow Project GitHub

- Flexibility ✓

- Execute on three different software environments
 - Conda
 - Docker
 - System

- Extensibility

- Create complicated Project Workflows

MLflow Project Documentation



Search

Documentation > MLflow Projects

Edit on GitHub



Quickstart

Tutorials and Examples

Concepts

MLflow Tracking

- MLflow Projects

Overview

+ Specifying Projects

+ Running Projects

Iterating Quickly

Building Multistep Workflows

MLflow Projects

An MLflow Project is a format for packaging data science code in a reusable and reproducible way, based primarily on conventions. In addition, the Projects component includes an API and command-line tools for running projects, making it possible to chain together projects into workflows.

Table of Contents

- [Overview](#)
- [Specifying Projects](#)
- [Running Projects](#)
- [Iterating Quickly](#)
- [Building Multistep Workflows](#)

<https://mlflow.org/docs/latest/projects.html>

MLflow Components

mlflow Tracking

Record and query experiments: code, data, config, and results

mlflow Projects

Package data science code in a format that enables reproducible runs on any platform

mlflow Models

Deploy machine learning models in diverse serving environments environments

new

mlflow Model Registry

Store, annotate and manage models in a central repository

[databricks.com
/mlflow](https://databricks.com/mlflow)



mlflow.org



github.com/mlflow



twitter.com/MLflow

MLflow Model Motivations



ML Frameworks

N x M
Combination of
Model support for
all Serving tools



Inference Code

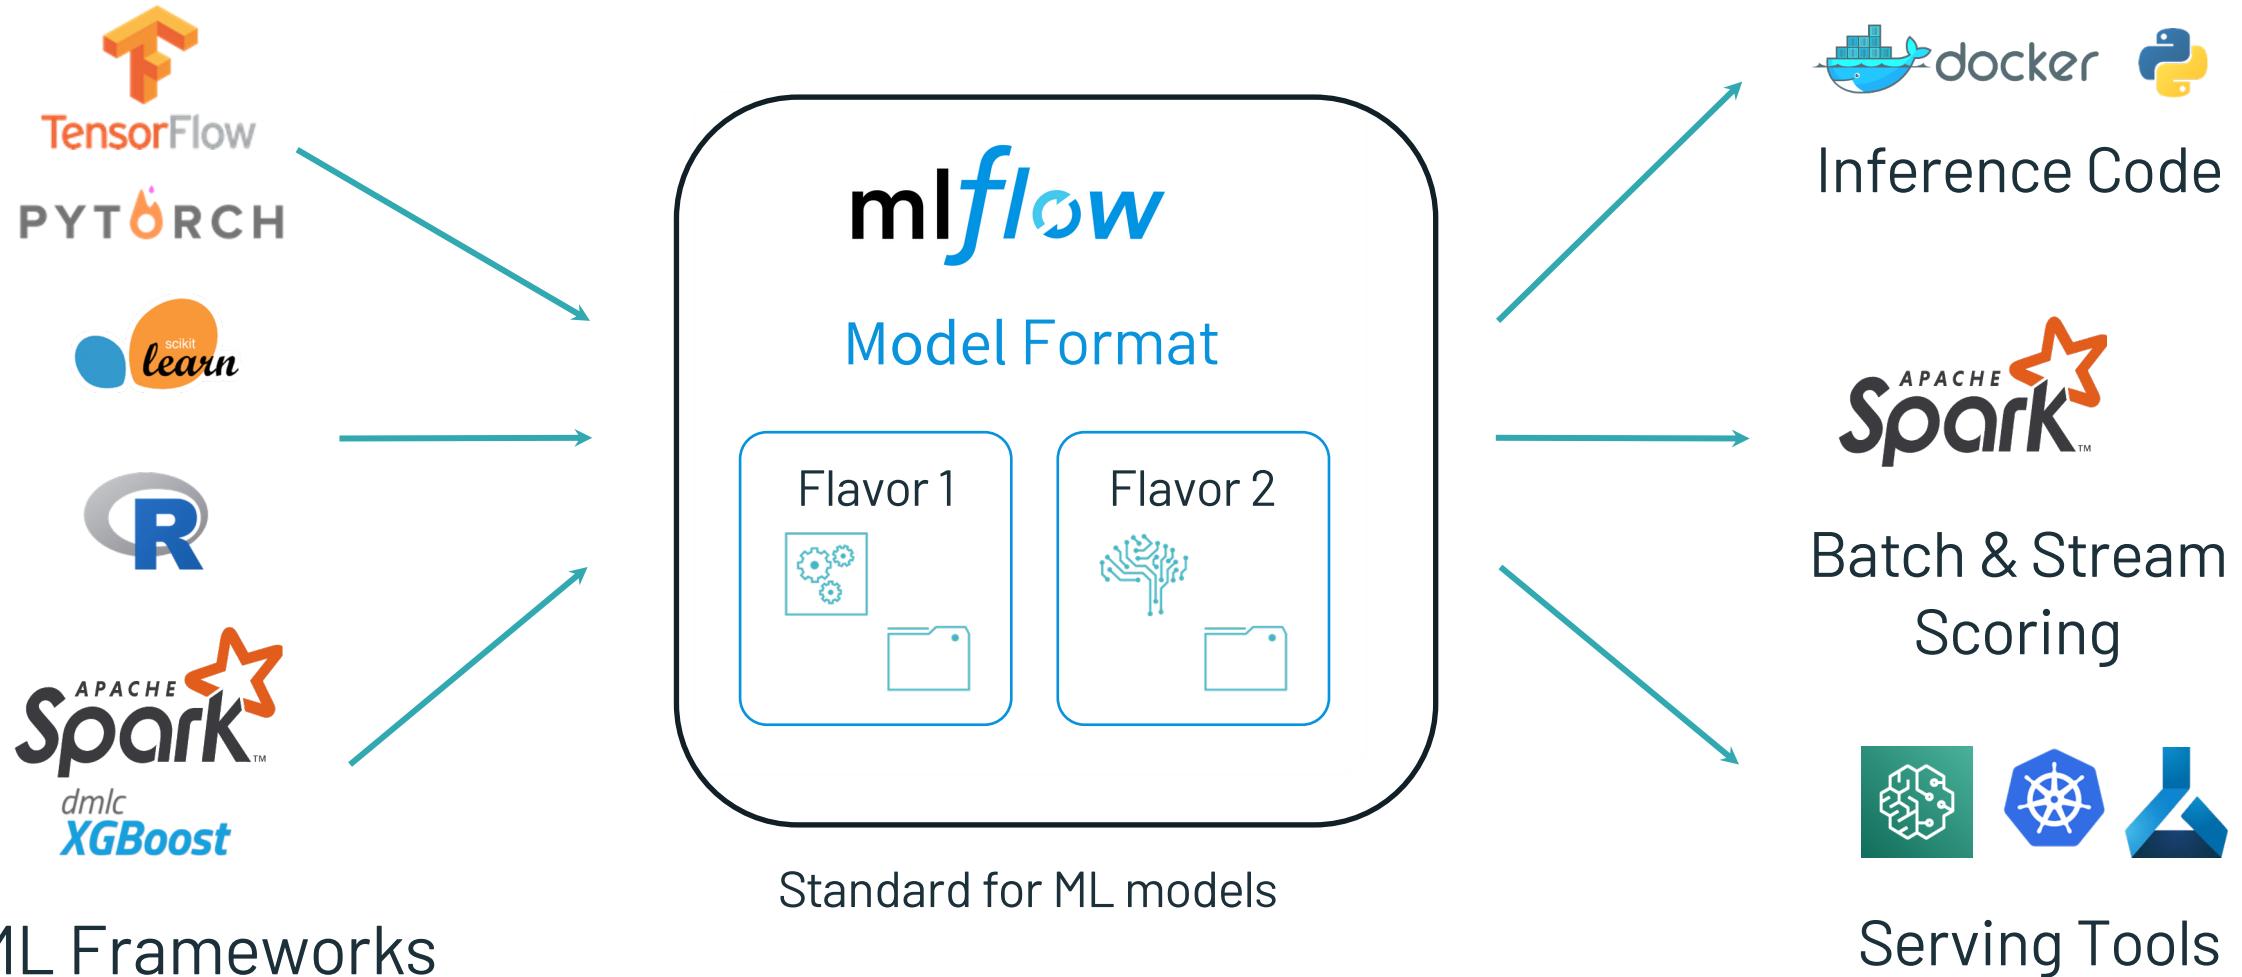


Batch & Stream Scoring



Serving Tools

MLflow Models

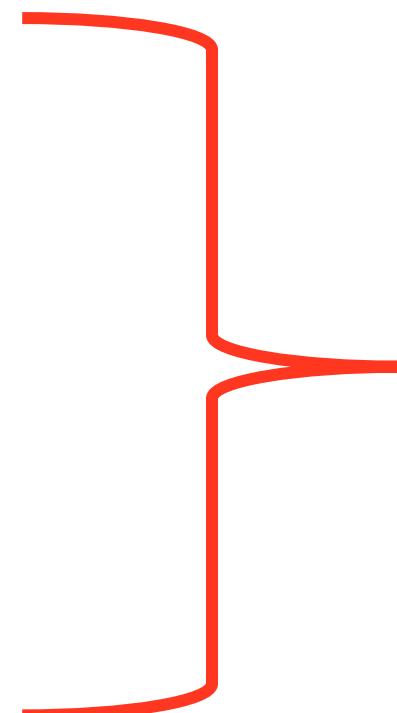


MLflow Model API Documentation

Built-In Model Flavors

MLflow provides several standard flavors that might be useful in your applications. Specifically, many of its deployment tools support these flavors, so you can export your own model in one of these flavors to benefit from all these tools:

- [Python Function \(python_function\)](#)
- [R Function \(crate\)](#)
- [H2O \(h2o\)](#)
- [Keras \(keras\)](#)
- [MLeap \(mleap\)](#)
- [PyTorch \(pytorch\)](#)
- [Scikit-learn \(sklearn\)](#)
- [Spark MLlib \(spark\)](#)
- [TensorFlow \(tensorflow\)](#)
- [ONNX \(onnx\)](#)
- [MXNet Gluon \(gluon\)](#)
- [XGBoost \(xgboost\)](#)
- [LightGBM \(lightgbm\)](#)
- [Spacy \(spaCy\)](#)
- [Fastai \(fastai\)](#)



Model flavor has generic methods:

- `log_model()`
- `save_model()`
- `load_model()`
- `add_model_flavor()`
- `auto_log() *`

* Only some model flavors

1. Example MLflow Model

```
mlflow.keras.log_model(...)
```

```
model
  └── MLmodel
  └── conda.yaml
  └── data
    └── keras_module.txt
    └── model.h5
```

1 directory, 4 files

```
artifact_path: model
flavors:
  keras:
    data: data
    keras_module: keras
    keras_version: 2.3.1
    python_function:
      data: data
      env: conda.yaml
      loader_module: mlflow.keras
      python_version: 3.7.5
...
run_id: 714d418027bc43bd8adc768f0f9ecb51
utc_time_created: '2020-08-25 19:04:53.887587'
```

Usable by tools that understand
Keras model format

Usable by any tool that can run
Python (Docker, Spark, etc!)

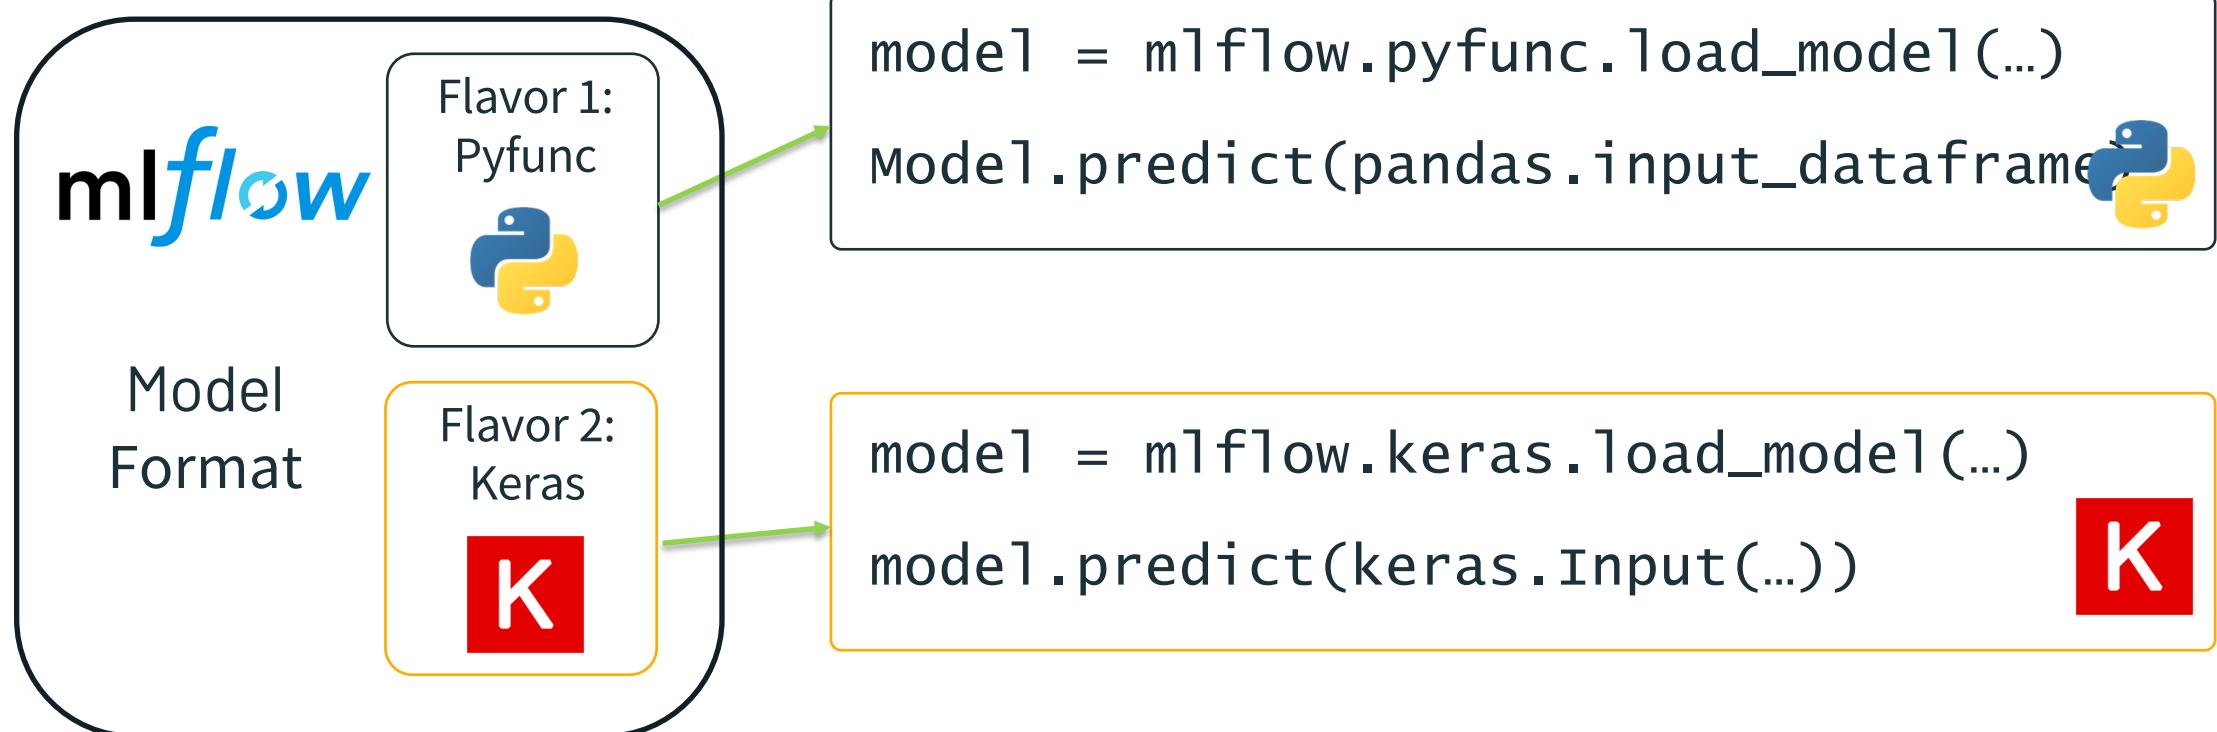
2. Example Conda.yaml

```
model
└── MLmodel
└── conda.yaml
└── data
    └── keras_module.txt
    └── model.h5
```

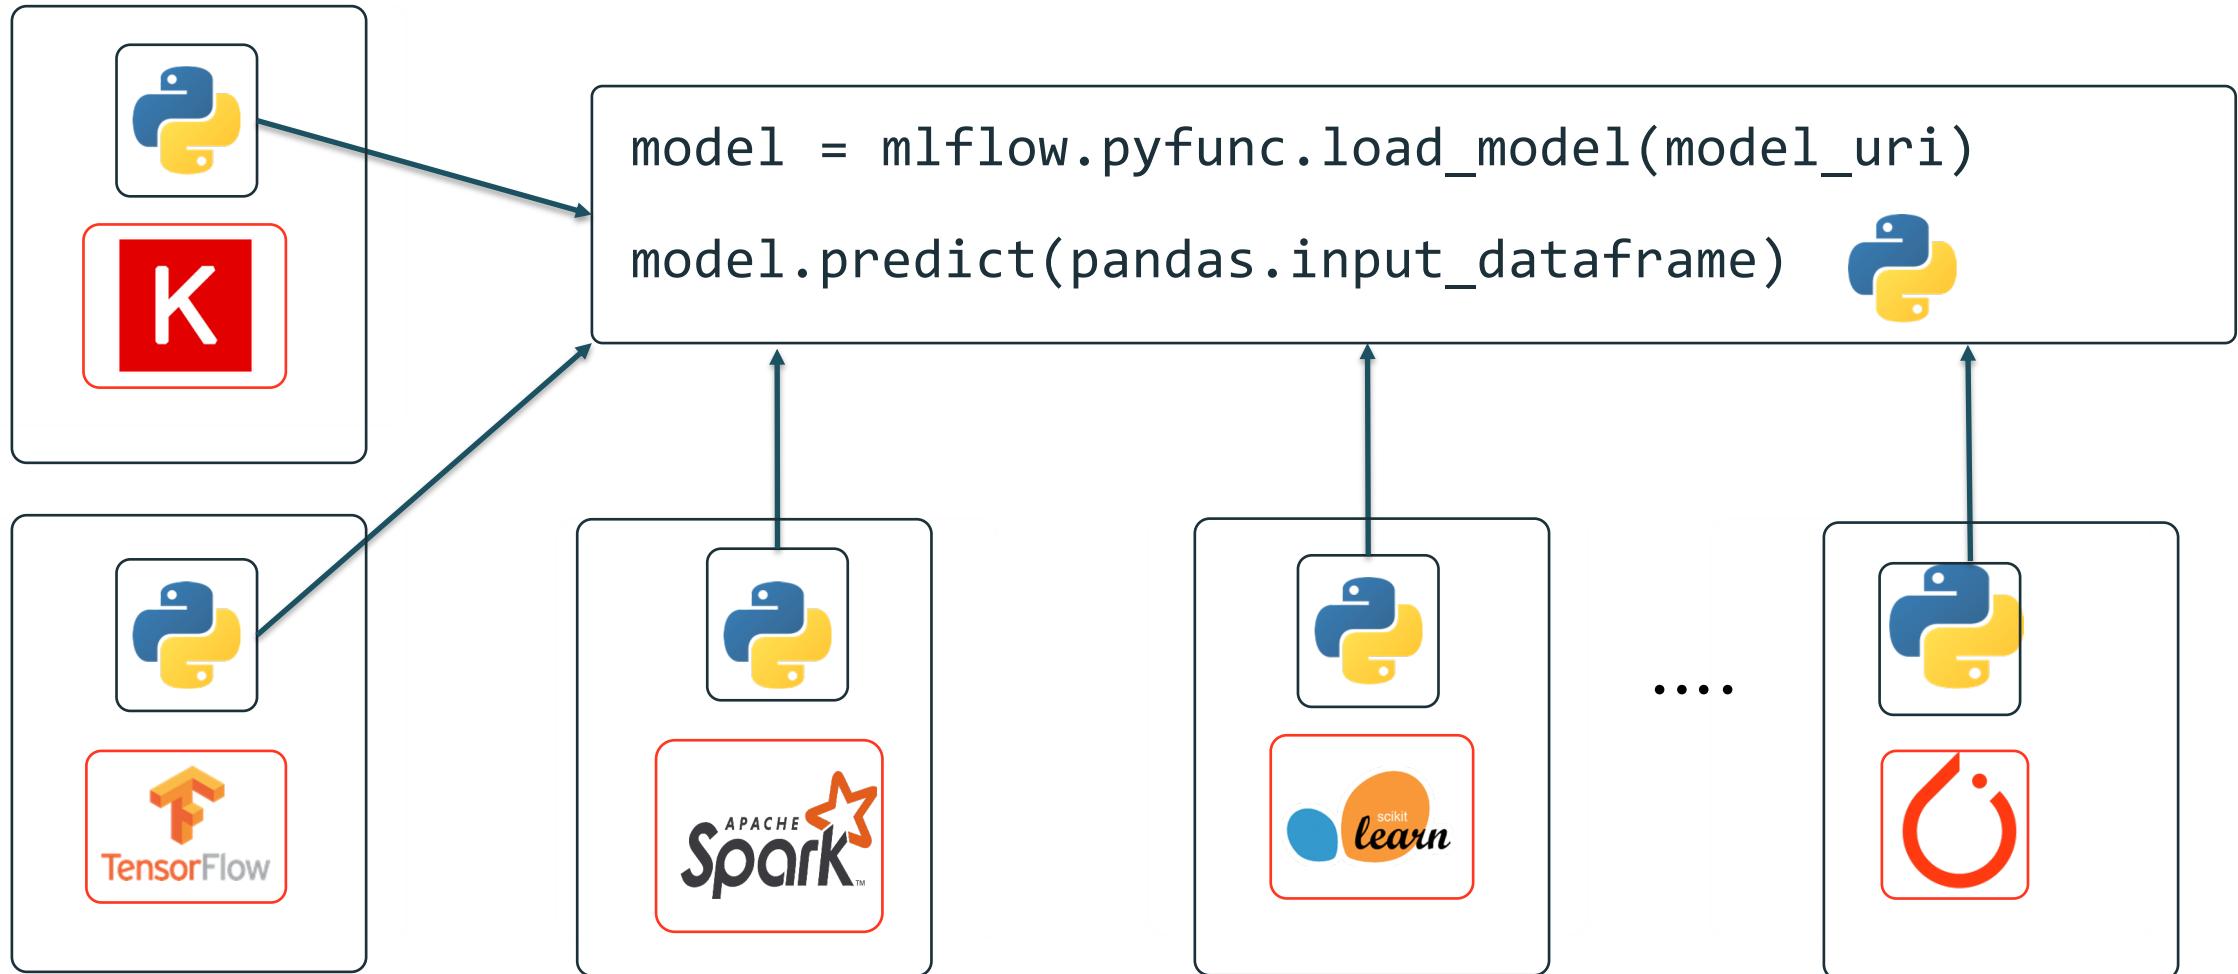
1 directory, 4 files

```
channels:
- defaults
- conda-forge
dependencies:
- python=3.7.5
- pip
- pip:
  - mlflow
  - keras==2.3.1
  - tensorflow==2.0.0
name: mlflow-env
```

Model Keras Flavor Example

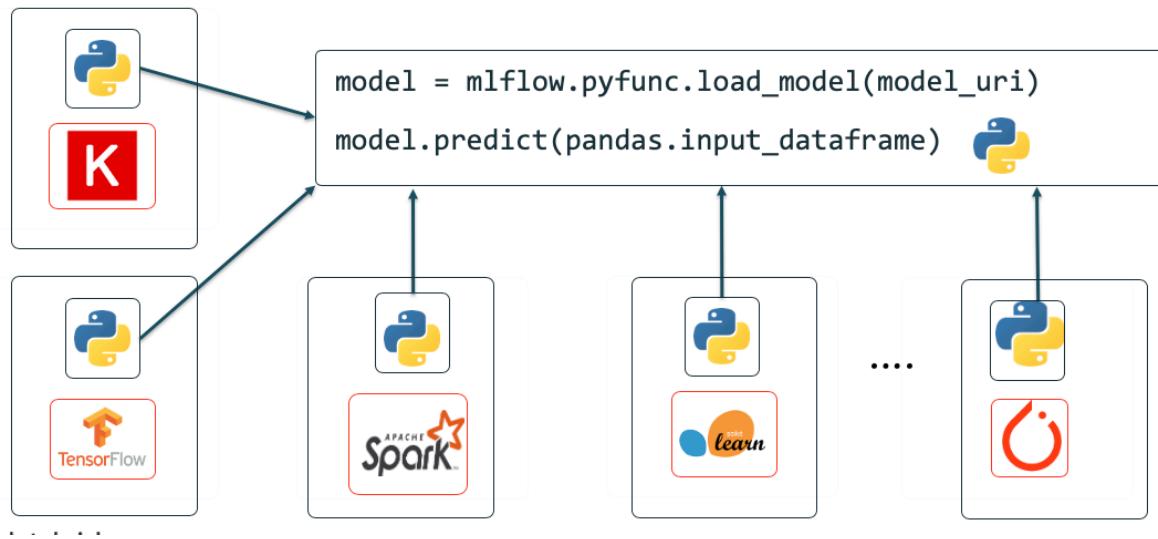


Model Flavors Example



MLflow PythonModel

- Serves as Generic Python model interface



- Use with Spark UDFs

```
predict = mlflow.pyfunc.spark_udf(spark, "runs://run_id/model")
df.withColumn("prediction", predict("name", "age")).show()
```

Recap: MLflow Models

Packaging format for ML Models

- Any directory with MLmodel file

Defines dependencies for reproducibility

- Conda environment can be specified in MLmodel configuration

Model creation and loading utilities

- mlflow.<model_flavor>.save_model(...) or log_model(...) or autolog()
- mlflow.<model_flavor>.load_model(...)

Deployment APIs

- CLI / Python / R / Java
- **mlflow models [OPTIONS] COMMAND [ARGS]...**
 - mlflow models serve [OPTIONS [ARGS]]
 - mlflow models predict [OPTIONS [ARGS] ...]
 - mlflow sagemaker -help
 - mlflow azureml --help

MLflow Components

mlflow Tracking

Record and query experiments: code, data, config, and results

mlflow Projects

Package data science code in a format that enables reproducible runs on any platform

mlflow Models

Deploy machine learning models in diverse serving environments

new

mlflow Model Registry

Store, annotate and manage models in a central repository

[databricks.com
/mlflow](https://databricks.com/mlflow)



mlflow.org



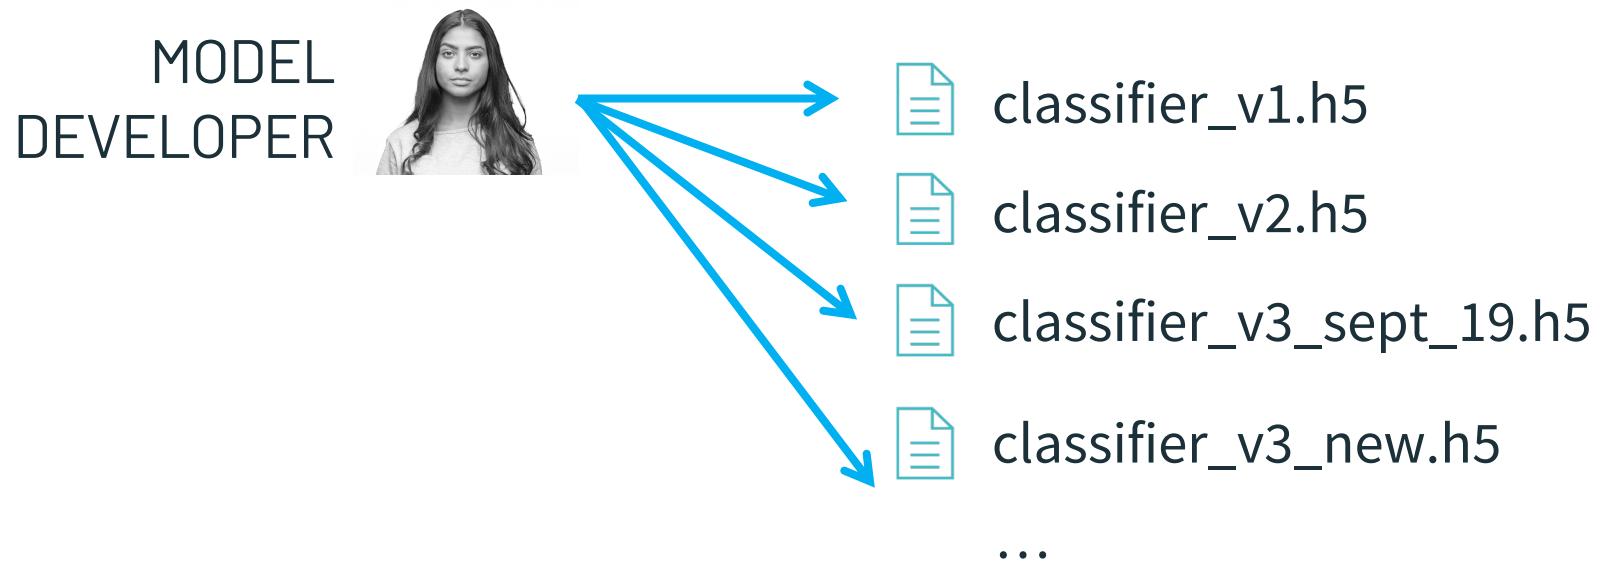
github.com/mlflow



twitter.com/MLflow

The Model Management Problem

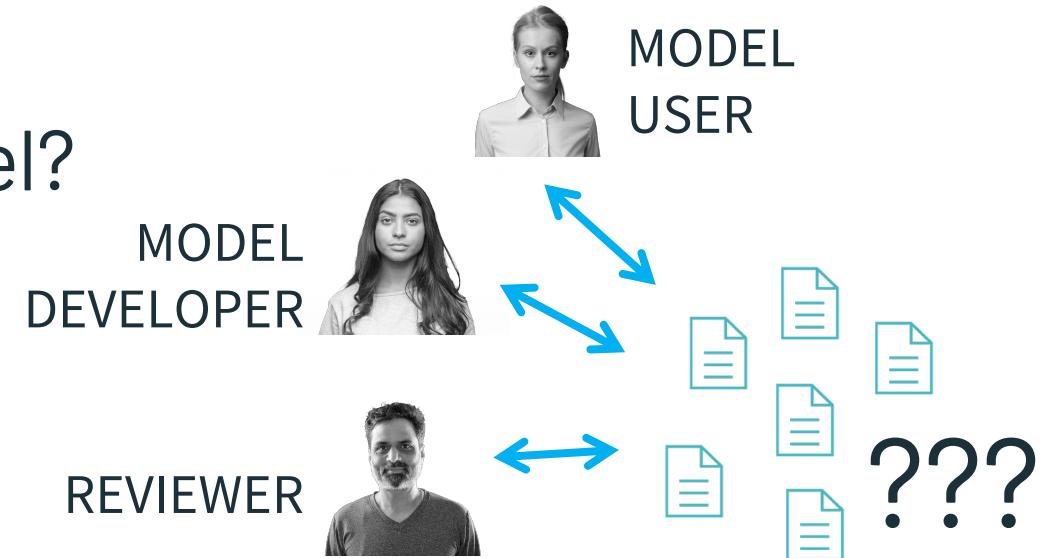
When you're working on one ML app alone, storing your models in files is manageable



The Model Management Problem

When you work in a large organization with many models, many data teams, management becomes a major challenge:

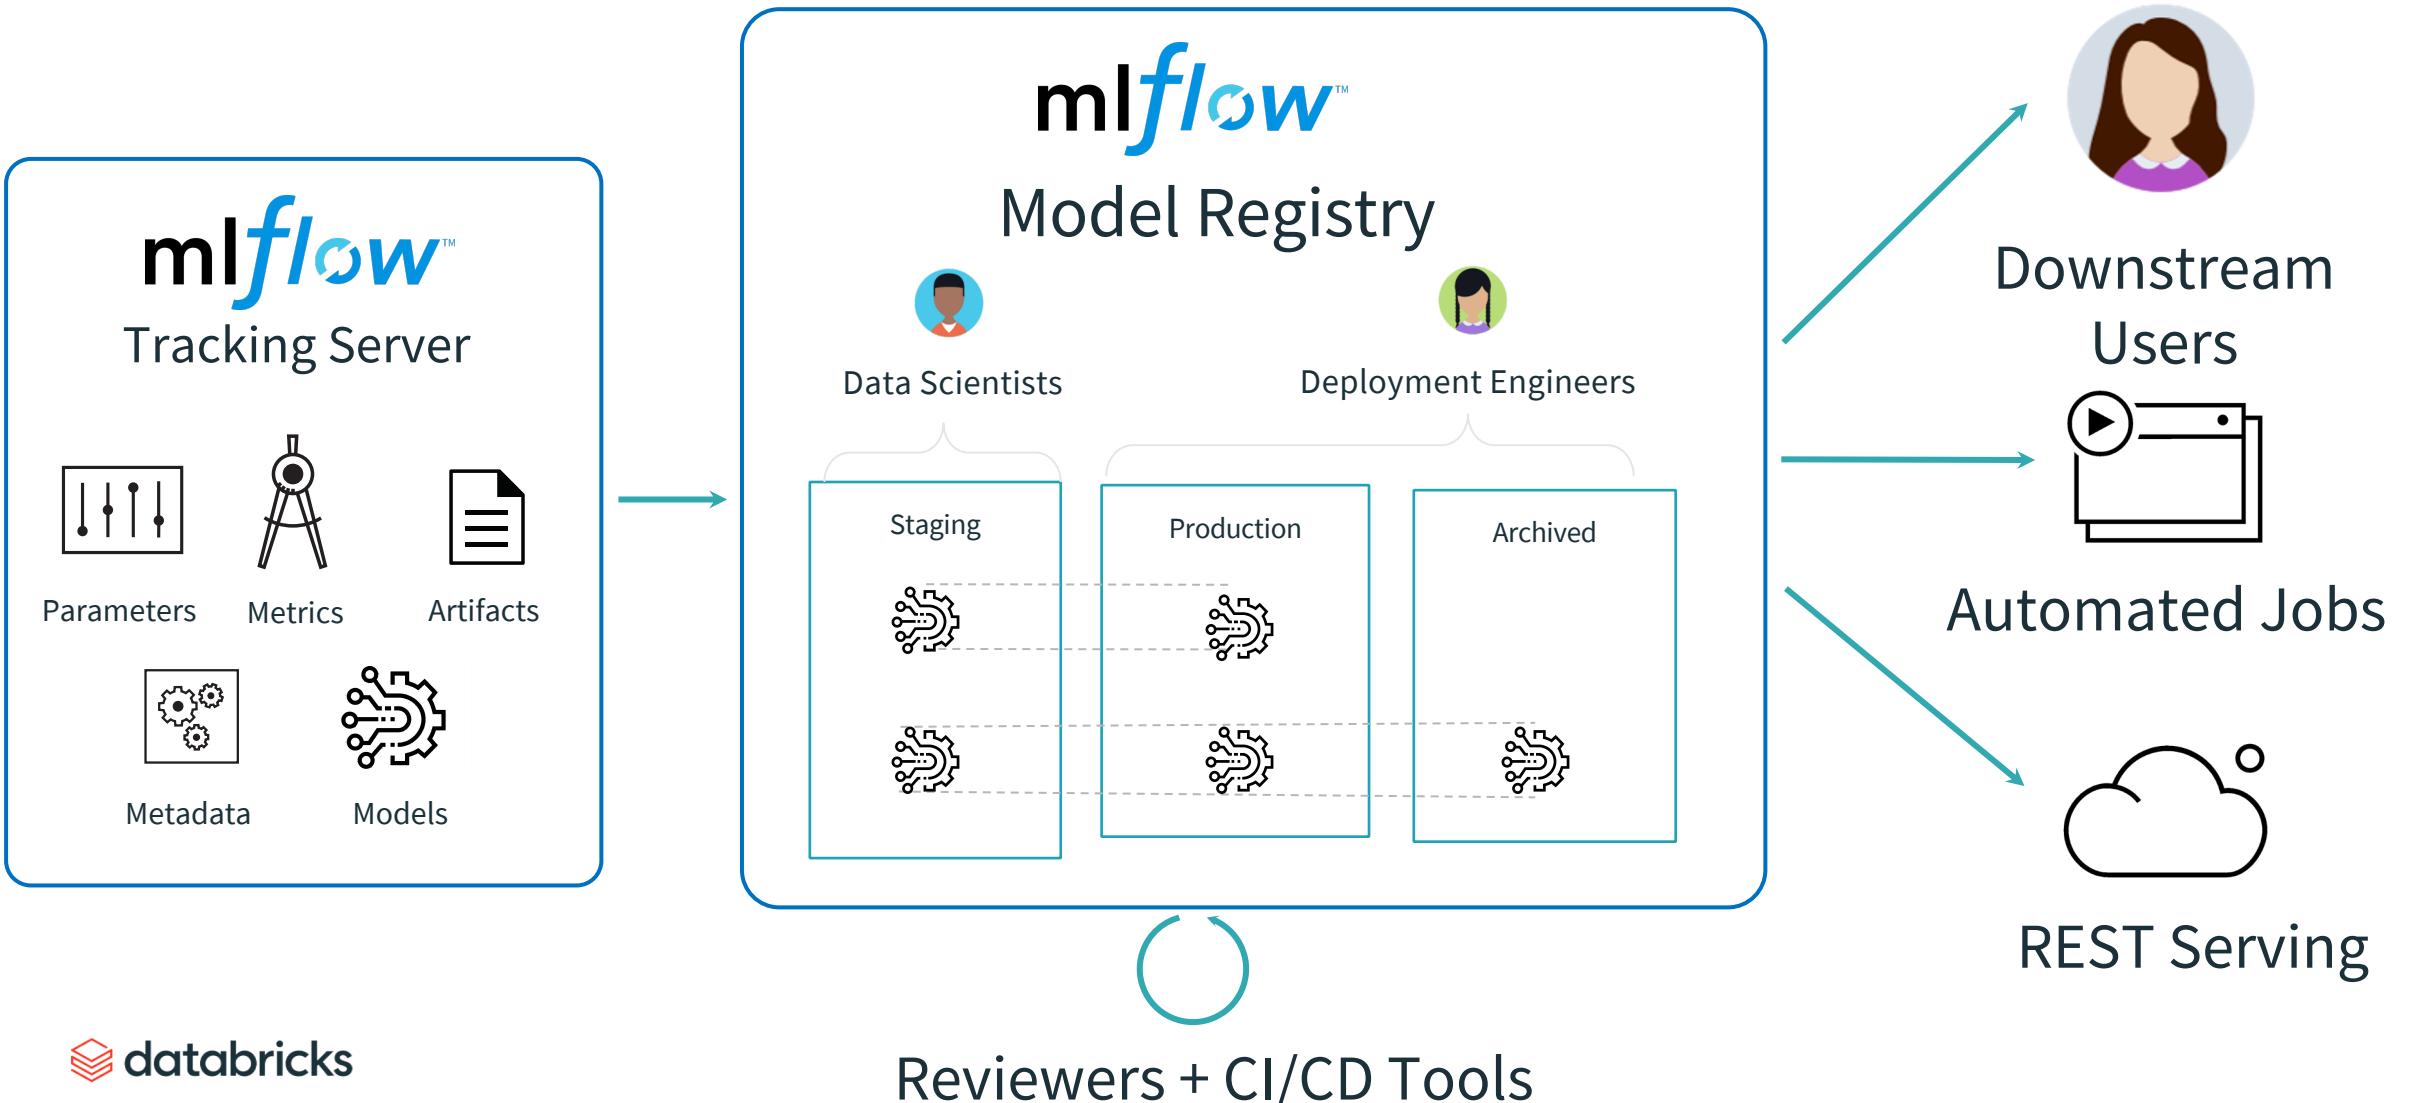
- Where can I find the best version of this model?
- How was this model trained?
- How can I track docs for each model?
- How can I review models?
- How can I integrate with CI/CD?





Model Registry

VISION: Centralized and collaborative model lifecycle management



MLflow Model Registry

- Repository of named, versioned models with controlled Access to Models
- Track each model's stage: none, staging, production, or archived
- Easily inspect a specific version and its run info
- Easily load a specific version
- Provides model description, lineage and activities

The screenshot shows the MLflow Model Registry interface for a registered model named "Airline_Delay_SparkML".

Header: Registered Models > Airline_Delay_SparkML

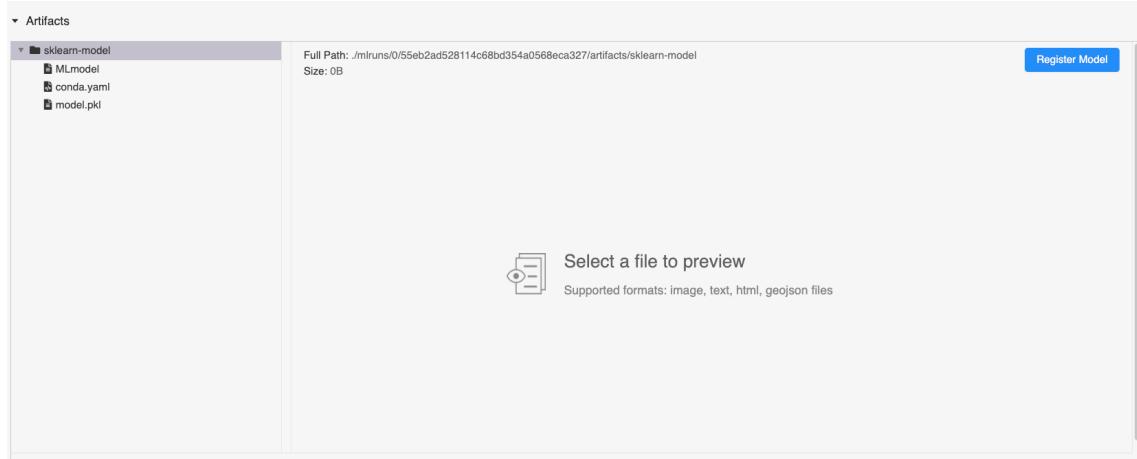
Created Time: 2019-10-10 15:20:29 **Last Modified:** 2019-10-14 12:17:04

Description: Predicts airline delays (in minutes) using the best Spark RF model from the AutoML Toolkit.

Versions: All Active(1)

Version	Registered at	Created by	Stage
Version 1	2019-10-10 15:20:30	clemens@demo.com	Archived
Version 2	2019-10-10 21:47:29	clemens@demo.com	Archived
Version 3	2019-10-10 23:39:43	clemens@demo.com	Production
Version 4	2019-10-11 09:55:29	clemens@demo.com	None
Version 5	2019-10-11 12:44:44	matei@demo.com	Staging

Model Registry Workflow UI



Artifacts

sklearn-model

- MLmodel
- conda.yaml
- model.pkl

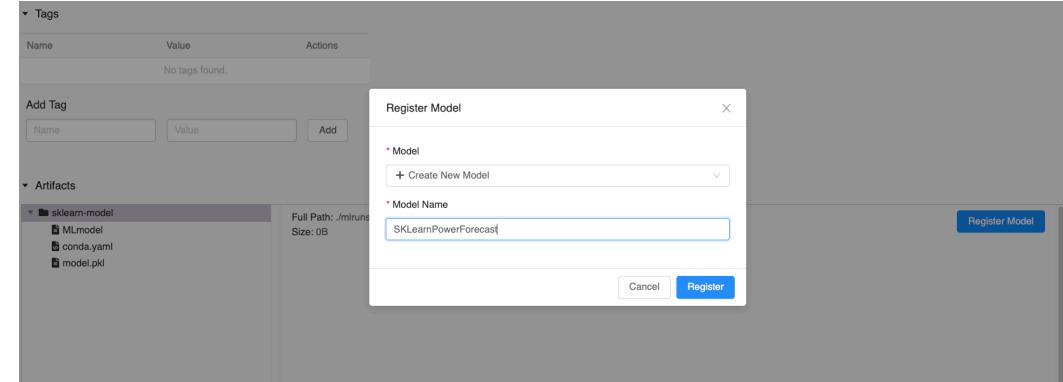
Full Path: ./mirluns/0/55eb2ad528114c68bd354a0568eca327/artifacts/sklearn-model
Size: 0B

Select a file to preview
Supported formats: image, text, html, geojson files

Register Model



MODEL
DEVELOPER



Tags

Name	Value	Actions
No tags found.		

Add Tag

Artifacts

sklearn-model

- MLmodel
- conda.yaml
- model.pkl

Full Path: ./mirluns/0/55eb2ad528114c68bd354a0568eca327/artifacts/sklearn-model
Size: 0B

Register Model

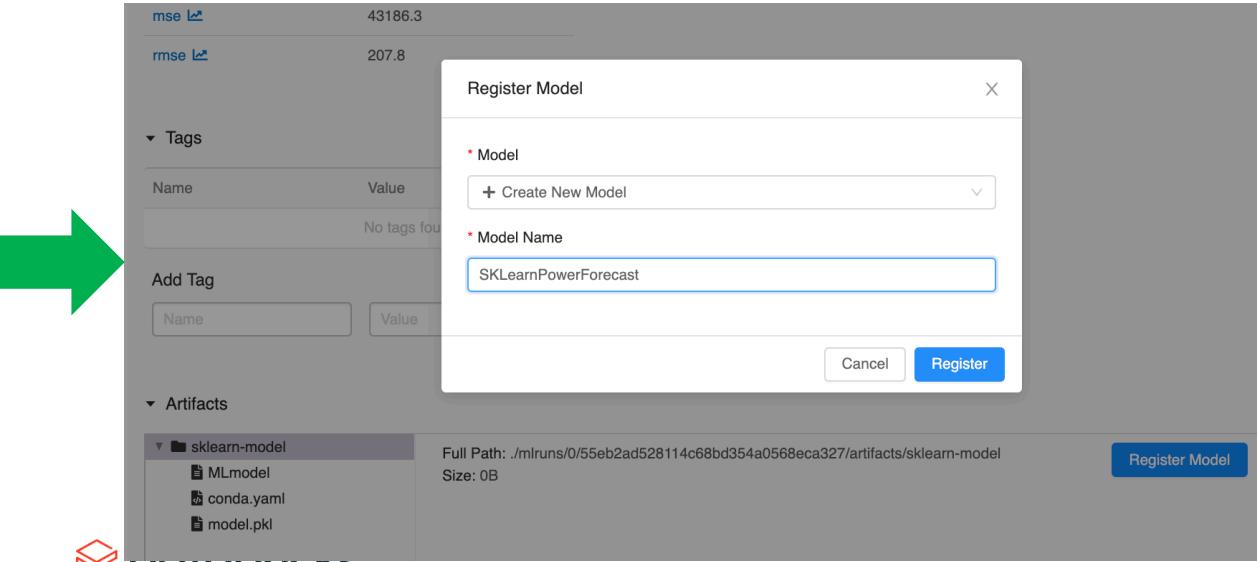
Model

+ Create New Model

Model Name

SKLearnPowerForecast

Cancel Register



mse ↗ 43186.3

rmse ↗ 207.8

Tags

Name	Value
No tags found.	

Add Tag

Name Value

Register Model

Model

+ Create New Model

Model Name

SKLearnPowerForecast

Cancel Register

Artifacts

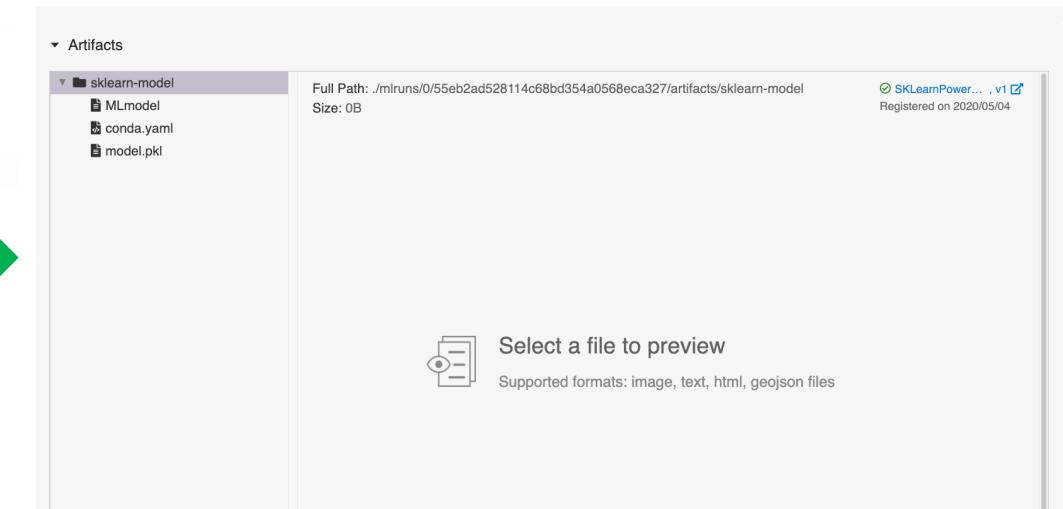
sklearn-model

- MLmodel
- conda.yaml
- model.pkl

Full Path: ./mirluns/0/55eb2ad528114c68bd354a0568eca327/artifacts/sklearn-model
Size: 0B

Select a file to preview
Supported formats: image, text, html, geojson files

Register Model



Artifacts

sklearn-model

- MLmodel
- conda.yaml
- model.pkl

Full Path: ./mirluns/0/55eb2ad528114c68bd354a0568eca327/artifacts/sklearn-model
Size: 0B

SKLearnPowerForecast, v1
Registered on 2020/05/04

Select a file to preview
Supported formats: image, text, html, geojson files

MLflow Model Registry

The MLflow Model Registry component is a centralized model store, set of APIs, and UI, to collaboratively manage the full lifecycle of an MLflow Model. It provides model lineage (which MLflow experiment and run produced the model), model versioning, stage transitions (for example from staging to production), and annotations.

Table of Contents

- Concepts
- Model Registry Workflows
 - UI Workflow
 - Registering a Model
 - Using the Model Registry
 - API Workflow
 - Adding an MLflow Model to the Model Registry
 - Fetching an MLflow Model from the Model Registry
 - Serving an MLflow Model from Model Registry
 - Adding or Updating an MLflow Model Descriptions
 - Renaming an MLflow Model
 - Transitioning an MLflow Model's Stage
 - Listing and Searching MLflow Models
 - Archiving an MLflow Model
 - Deleting MLflow Models



`mlflow.register_model(model_uri, name)` [\[source\]](#) 

`mlflow.get_registry_uri()` [\[source\]](#)

`mlflow.set_registry_uri(uri)` [\[source\]](#)

Model Registry Workflow UI

This screenshot shows the mlflow Model Registry UI. It displays a model version named 'SKLearnPowerForecast > Version 1'. The page includes details such as 'Registered At: 2020-05-04 11:38:47', 'Creator:', 'Stage: None', and 'Source Run: Random Forest Regressor: Power Forecasting Model'. A large green arrow points from this screen to the 'MODEL REVIEWER' section.



MODEL
REVIEWER

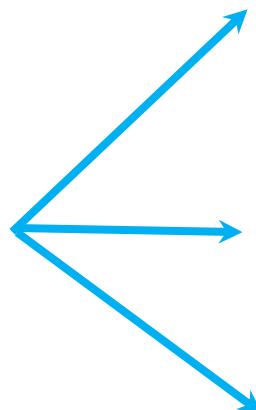
This screenshot shows the mlflow Model Registry UI. It displays a model version named 'SKLearnPowerForecast > Version 1'. The 'Stage' dropdown is set to 'None'. Below it, there are three buttons for transitioning the model: 'Transition to → Staging' (orange), 'Transition to → Production' (green), and 'Transition to → Archived' (grey). A large green arrow points from the 'MODEL REVIEWER' section to this screen.



DOWNSTREAM
USERS

This screenshot shows the mlflow Model Registry UI. It displays a list of model versions for 'SKLearnPowerForecast'. The table includes columns for 'Version', 'Registered at', 'Created by', and 'Stage'. Three versions are listed: Version 1 (Staging), Version 2 (Staging), and Version 3 (Production). A large green arrow points from the 'DOWNSTREAM USERS' section to this screen.

Version	Registered at	Created by	Stage
Version 1	2020-05-04 11:38:47		Staging
Version 2	2020-05-04 11:41:37		Staging
Version 3	2020-05-04 11:42:07		Production



AUTOMATED JOBS



REST SERVING

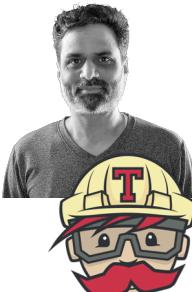
Model Registry Workflow API

```
mlflow.register_model(model_uri, "WeatherForecastModel")  
  
mlflow.sklearn.log_model(model,  
    artifact_path="sklearn_model",  
    registered_model_name= "WeatherForecastModel")
```

MODEL
DEVELOPER



REVIEWERS,
CI/CD TOOLS



```
client = mlflow.tracking.Mlflowclient()  
client.transition_model_version_stage(name="WeatherForecastModel",  
    version=5,  
    stage="Production")
```

```
model_uri= "models:/{{model_name}}/production".format(  
    model_name="WeatherForecastModel")  
model_prod = mlflow.sklearn.load_model(model_uri)  
model_prod.predict(data)
```

DOWNSTREAM
USERS



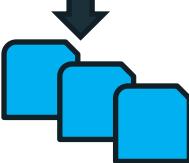
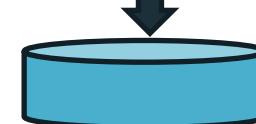
AUTOMATED JOBS



REST SERVING



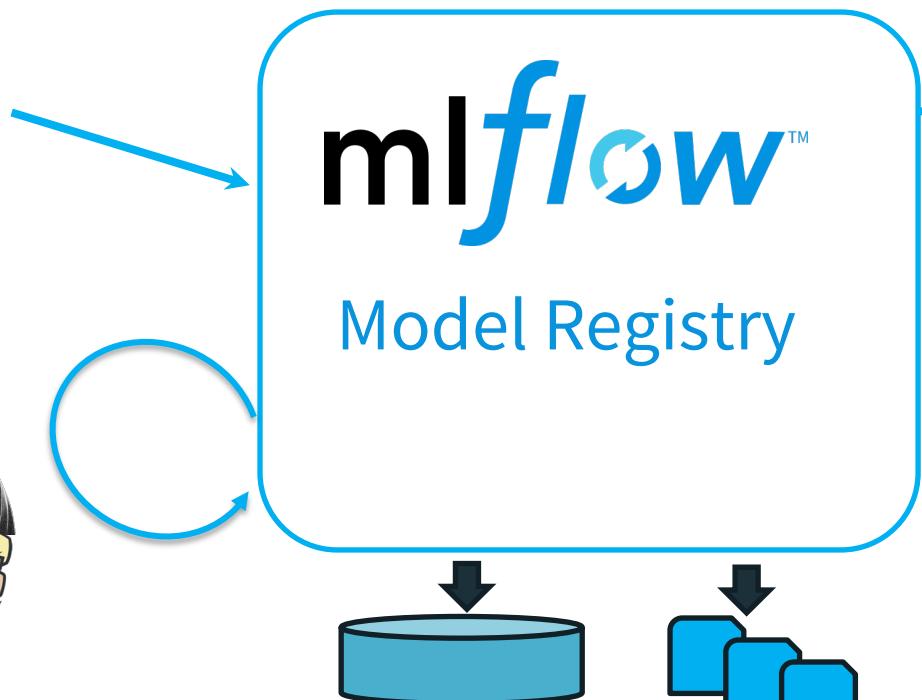
mlflow™
Model Registry



Model Registry Workflow API

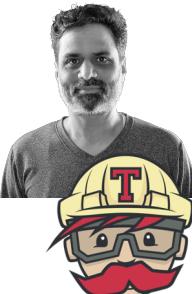
```
mlflow.register_model(model_uri, "WeatherForecastModel")  
  
mlflow.sklearn.log_model(model,  
    artifact_path="sklearn_model",  
    registered_model_name= "WeatherForecastModel")
```

MODEL
DEVELOPER

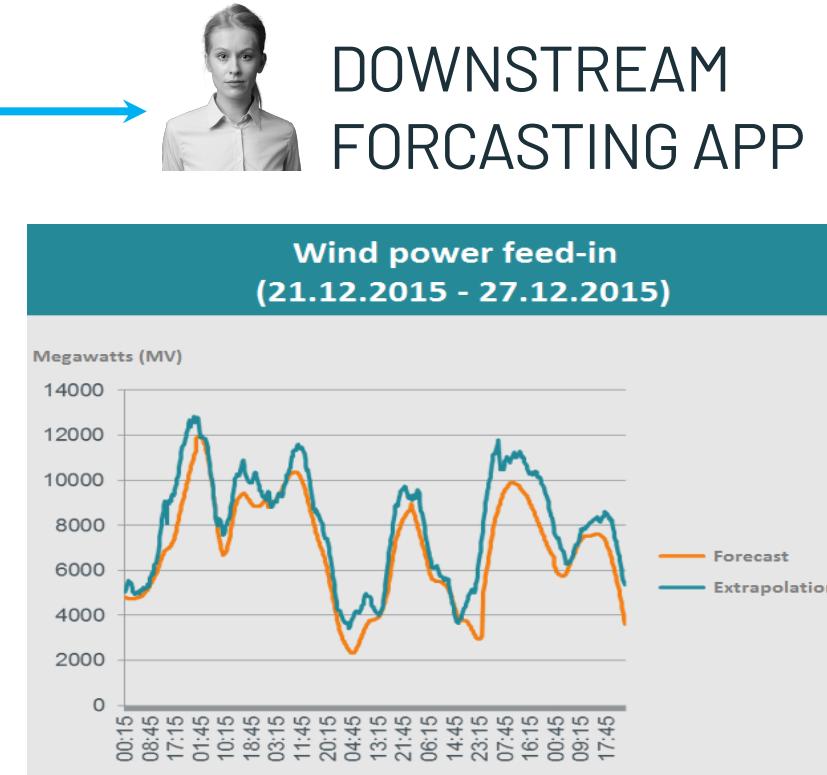


```
model_uri = "models:/{{model_name}}/production".format(  
    model_name="WeatherForecastModel")  
model_prod = mlflow.pyfunc.load_model(model_uri)  
model_prod.predict(data)
```

REVIEWERS,
CI/CD TOOLS



```
client = mlflow.tracking.MlflowClient()  
client.transition_model_version_stage(name="WeatherForecastModel",  
    version=5,  
    stage="Production")
```



MLflow Model Registry

The MLflow Model Registry component is a centralized model store, set of APIs, and UI, to collaboratively manage the full lifecycle of an MLflow Model. It provides model lineage (which MLflow experiment and run produced the model), model versioning, stage transitions (for example from staging to production), and annotations.

Table of Contents

- [Concepts](#)
- [Model Registry Workflows](#)

- [UI Workflow](#)
 - [Registering a Model](#)
 - [Using the Model Registry](#)

- [API Workflow](#)
 - [Adding an MLflow Model to the Model Registry](#)
 - [Fetching an MLflow Model from the Model Registry](#)
 - [Serving an MLflow Model from Model Registry](#)
 - [Adding or Updating an MLflow Model Descriptions](#)
 - [Renaming an MLflow Model](#)
 - [Transitioning an MLflow Model's Stage](#)
 - [Listing and Searching MLflow Models](#)
 - [Archiving an MLflow Model](#)
 - [Deleting MLflow Models](#)

Model Registry CRUD Operations MLflowClient()

`create_model_version(name, source, run_id, tags=None, run_link=None, description=None)` [\[source\]](#)

`create_registered_model(name, tags=None, description=None)` [\[source\]](#)

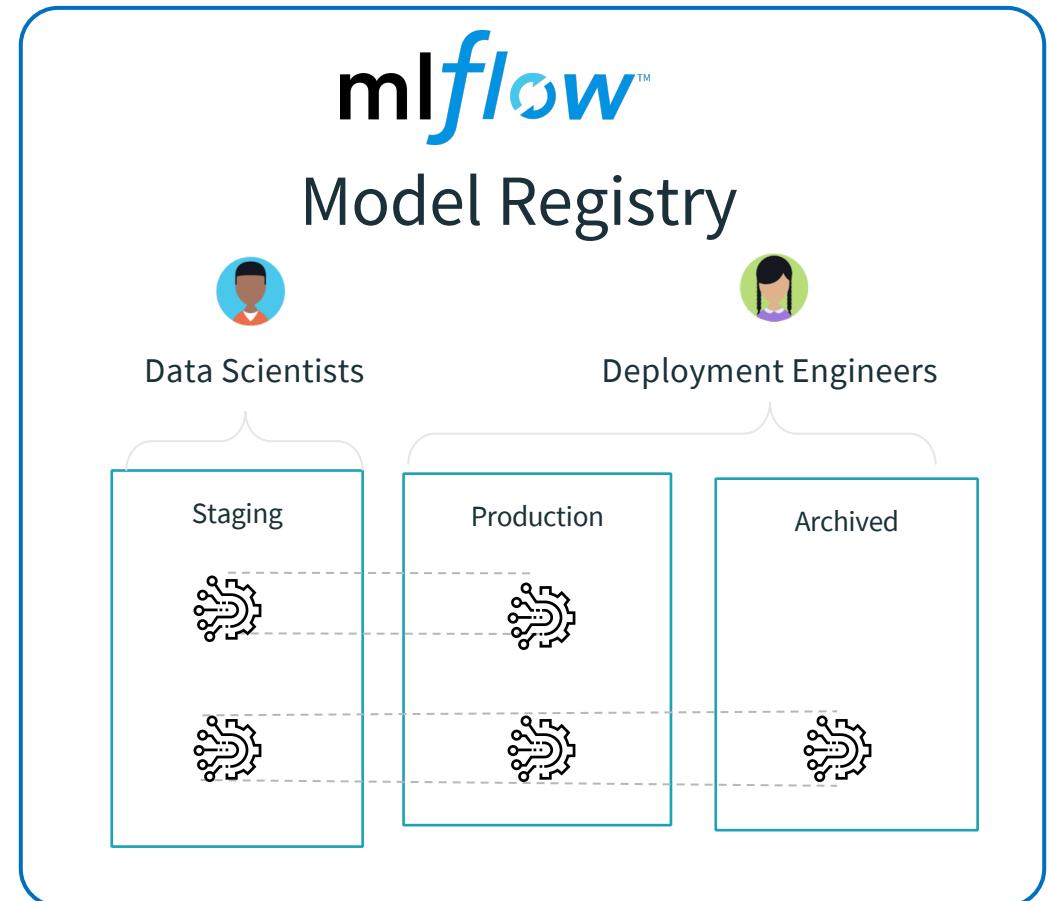
`delete_model_version(name, version)` [\[source\]](#)

`get_latest_versions(name, stages=None)` [\[source\]](#)

`transition_model_version_stage(name, version, stage, archive_existing_versions=False)` [\[source\]](#)

MLflow Model Registry Recap

- **Central Repository:** Unique named registered models for discovery across data teams
- **Model Registry Workflow:** Provides UI and API for registry operations
- **Model Versioning:** Allow multiple versions of model in different stages
- **Model Stages:** Allow stage transition: none, staging, production, or archived
- **CI/CD Integration:** Easily load a specific version for testing and inspection
- **Model Lineage:** Provides model description, lineage and activities



Recap of all parts: What Did We Talk About?



- Modular Components greatly simplify the ML lifecycle
- Easy to install & Great Developer experience
- Develop & Deploy locally; track locally or remotely
- Available APIs: Python, Java & R (Soon Scala)
- REST APIs and CLI tools
- Visualize experiments and compare runs
- Centrally register and manage model lifecycle

Learning More About MLflow & Get Involved!

- `pip install mlflow` – to get started
- Find docs & examples at mlflow.org
- Peruse code and contribute at [MLflow Github](https://github.com/mlflow/mlflow)
- Join the [Slack channel](#)
- More [MLflow tutorials](#)

MLflow Tutorial

Tutorials: <https://github.com/dmatrix/ds4g-workshop>

Thank you! 😊

Q & A

jules@databricks.com
@2twitme

<https://www.linkedin.com/in/dmatrix/>